

ITP 30002 Operating System

Limited Direct Execution and Context Switching

OSTEP Chapter 6

Shin Hong

Mechanisms for Process

2

- time-sharing
 - run one process for a while, and then switch to another one
 - issue
 - managing control
 - obtaining performance
- limited direct execution
 - run an application program directly on the CPU with some restriction
 - restriction
 - restricted memory accesses (H/W manipulation, resource allocation, etc.)
 - restricted instruction



Process
--
ITP 30002
Operating System

2021-03-20

Dual Mode Operation: User Mode & Kernel Mode

3

- Most contemporary processors provide dual mode operation
- User mode
 - for application program execution
 - restriction is enforced
 - a trap occurs when a process executes a restricted instruction under user mode
- Kernel model
 - for kernel execution
 - all privileged operations can be executed
- What if an application program needs to execute privileged operations?

Process
--
ITP 30002
Operating System

2021-03-20

System Call

4

- A way of an application program to call a kernel to get a system service
 - not possible to do this with the procedure call mechanism
- to execute a system call, a program must execute a trap instruction
 - the operation of a trap instruction
 - change the mode into the kernel mode
 - store the current PC at a kernel stack
 - jump to a predefined program location for handling the trap
 - trap table
 - trap handler
 - each system call is identified by its unique number (i.e. system-call number)

Process
--
ITP 30002
Operating System

2021-03-20

Example of System Call Workflow

5

OS @ run (kernel mode)	Hardware	Program (user mode)
Create entry for process list Allocate memory for program Load program into memory Setup user stack with argv Fill kernel stack with reg/PC return-from-trap		
	restore regs (from kernel stack) move to user mode jump to main	Run main() ... Call system call trap into OS
	save regs (to kernel stack) move to kernel mode jump to trap handler	
Handle trap Do work of syscall return-from-trap		
	restore regs (from kernel stack) move to user mode jump to PC after trap	... return from main trap (via <code>exit()</code>)
Free memory of process Remove from process list		

Process
--
ITP 30002
Operating System

2021-03-20

Switching Between Processes

6

- How can OS regain control of the CPU when it is given to an application program?
- Natural chance: blocked operation
- Periodic scheduling
 - Cooperative approach
 - blocked operation
 - Preemptive scheduling approach
 - exploit a timer interrupt and its interrupt handler
 - requires HW support

Process
--
ITP 30002
Operating System

2021-03-20

Preemptive Scheduling

7

- scheduler: a kernel module that determines which process to dispatch at a chance (e.g., timer interrupt)
- context switch: a kernel module that is executed to switch processes running on a CPU
 - registered as a timer interrupt handler
 - steps
 - store the process status of a currently-running process to memory
 - CPU states
 - find the stored status of the next process from the memory
 - restore the stored status at the CPU
 - return the control back to the application program

Example

OS @ boot (kernel mode)	Hardware	
initialize trap table	remember addresses of... syscall handler timer handler	
start interrupt timer	start timer interrupt CPU in X ms	
OS @ run (kernel mode)	Hardware	Program (user mode)
		Process A
		...
	timer interrupt save regs(A) → k-stack(A) move to kernel mode jump to trap handler	
Handle the trap Call <code>switch()</code> routine save regs(A) → <code>proc_t(A)</code> restore regs(B) ← <code>proc_t(B)</code> switch to k-stack(B) return-from-trap (into B)		
	restore regs(B) ← k-stack(B) move to user mode jump to B's PC	
		Process B
		...

The xv6 Context Switch Code

9

```
1  # void swtch(struct context **old, struct context *new);
2  #
3  # Save current register context in old
4  # and then load register context from new.
5  .globl swtch
6  swtch:
7      # Save old registers
8      movl 4(%esp), %eax # put old ptr into eax
9      popl 0(%eax)      # save the old IP
10     movl %esp, 4(%eax) # and stack
11     movl %ebx, 8(%eax) # and other registers
12     movl %ecx, 12(%eax)
13     movl %edx, 16(%eax)
14     movl %esi, 20(%eax)
15     movl %edi, 24(%eax)
16     movl %ebp, 28(%eax)
17
18     # Load new registers
19     movl 4(%esp), %eax # put new ptr into eax
20     movl 28(%eax), %ebp # restore other registers
21     movl 24(%eax), %edi
22     movl 20(%eax), %esi
23     movl 16(%eax), %edx
24     movl 12(%eax), %ecx
25     movl 8(%eax), %ebx
26     movl 4(%eax), %esp # stack is switched here
27     pushl 0(%eax)      # return addr put in place
28     ret               # finally return into new ctxt
```

Process
--
ITP 30002
Operating System

2021-03-20