

ITP 30002 Operating System

# Paging

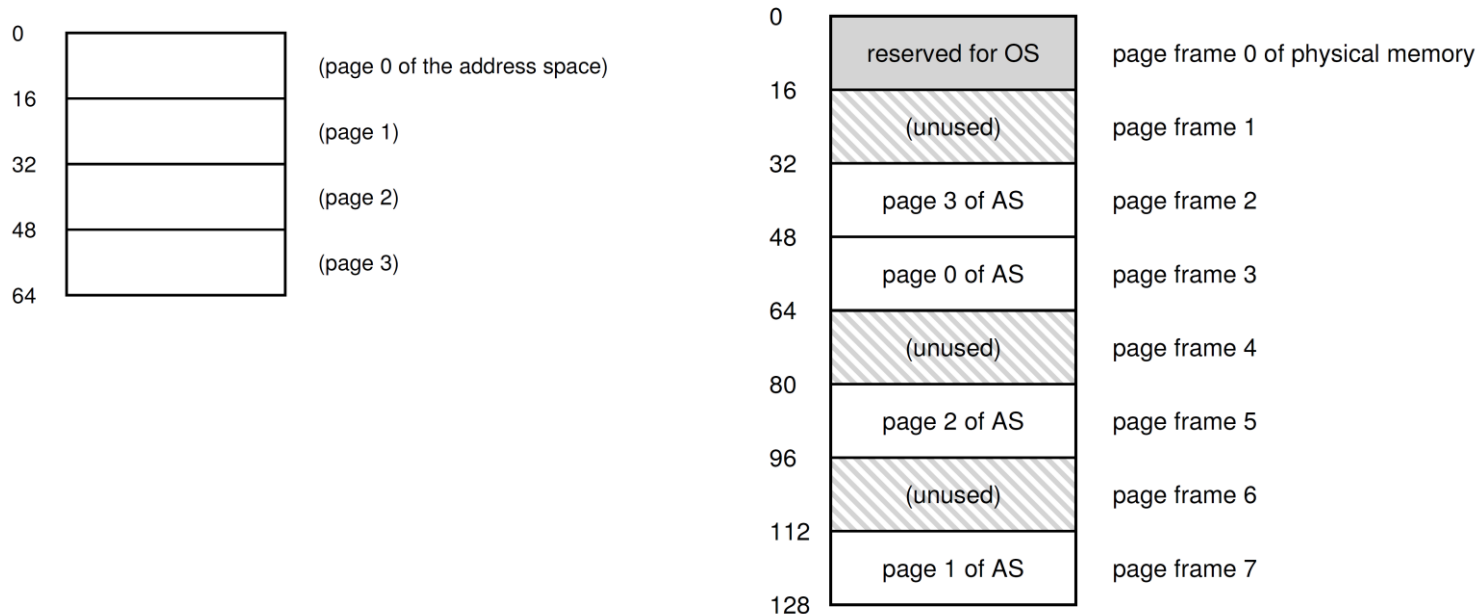
OSTEP Chapters 18 and 19

Shin Hong

# Paging

2

- allocate a memory resource to a process as a set of small fixed-size pieces
  - divide an address space into **pages** and divide a physical memory space into **frames** such that the sizes of a page and a frame are the same
  - can resolve both internal and external fragmentation problems if the maintenance overhead is manageable
- ex. 16-byte pages/frames



Paging

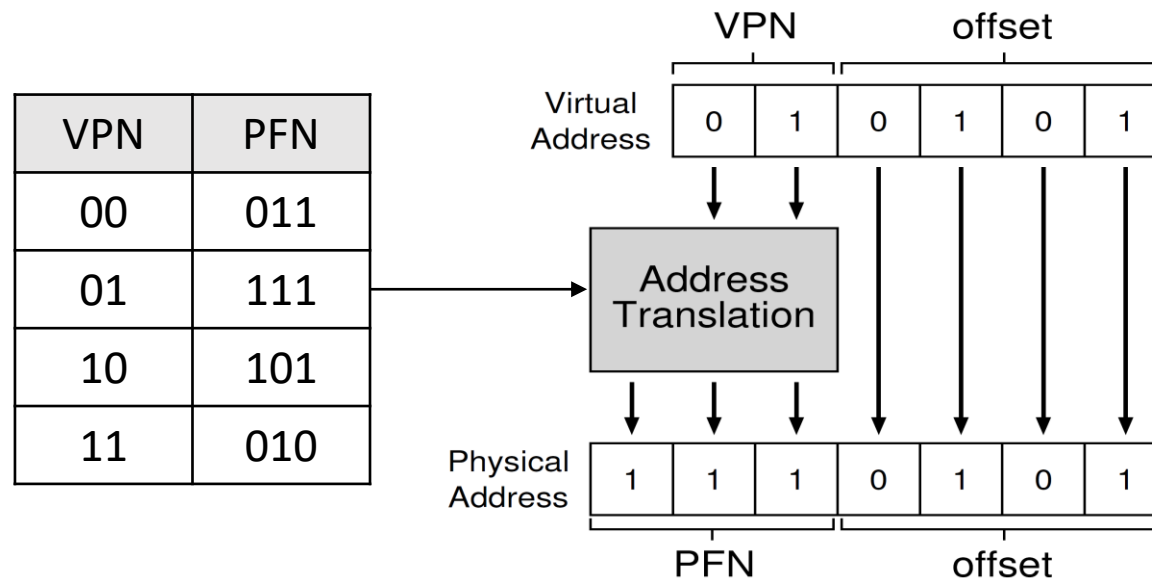
ITP 30002  
Operating System

2021-04-14

# Page Table and Address Translation

3

- there must be a page table which records to which frame a page of a process is assigned
  - structured as per-process page table, or inverted page table
- to translate a virtual address, we have to split it into the VPN and the offset, and replace VPN with the corresponding PFN
  - ex. a 64-byte address space in a 128-byte physical memory



Paging


ITP 30002  
Operating System

2021-04-14

# Page Table Entry

4

- valid bit: whether or not the page is in use
  - permission bits: read-only, read-write, supermode-only, etc.
  - present bit: whether the page is in physical memory or on disk
  - dirty bit: whether the page was modified since it was brought into memory
  - reference bit: whether the page has been recently accessed
- 
- Example of 32-bit x86



Paging

ITP 30002  
Operating System

2021-04-14

# Page Table Size

5

- page tables take a large amount of memory
  - example
    - 4 KB page/frame in a 32-bit address space and a 4 GB memory
    - 20-bits VPN,  $2^{20}$  page table entries
    - 4 bytes for each entry, thus, 4 MB ( $=2^{22}$  bytes) for a page table
    - 400 MB for 100 processes
- a page table should be resided in a main memory since MMU cannot accommodate a full page table
  - every translation step requires at least one memory access

Paging

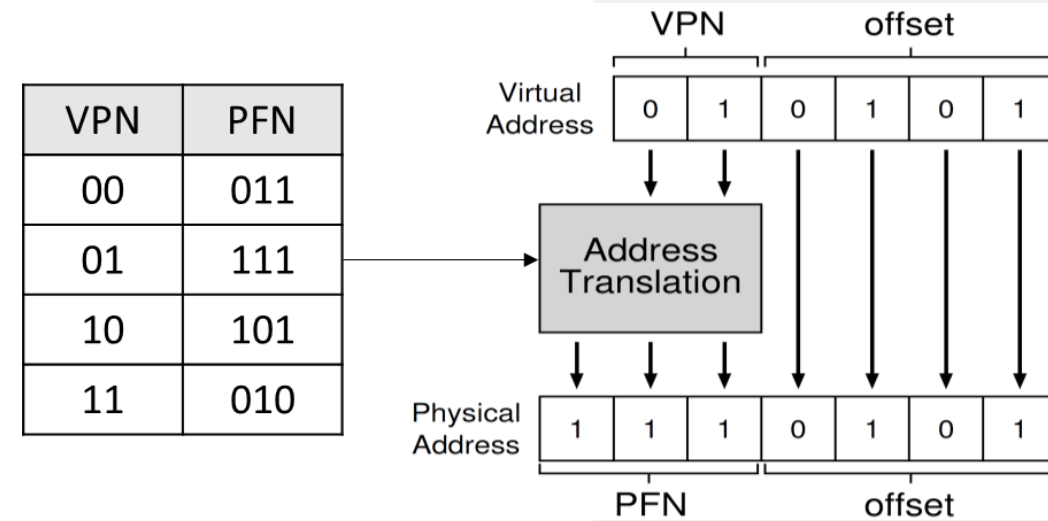
ITP 30002  
Operating System

2021-04-14

# Accessing Memory with Paging

6

```
1 // Extract the VPN from the virtual address
2 VPN = (VirtualAddress & VPN_MASK) >> SHIFT
3
4 // Form the address of the page-table entry (PTE)
5 PTEAddr = PTBR + (VPN * sizeof(PTE))
6
7 // Fetch the PTE
8 PTE = AccessMemory(PTEAddr)
9
10 // Check if process can access the page
11 if (PTE.Valid == False)
12     RaiseException(SEGMENTATION_FAULT)
13 else if (CanAccess(PTE.ProtectBits) == False)
14     RaiseException(PROTECTION_FAULT)
15 else
16     // Access is OK: form physical address and fetch it
17     offset = VirtualAddress & OFFSET_MASK
18     PhysAddr = (PTE.PFN << PFN_SHIFT) | offset
19     Register = AccessMemory(PhysAddr)
```



Paging

ITP 30002  
Operating System

2021-04-14

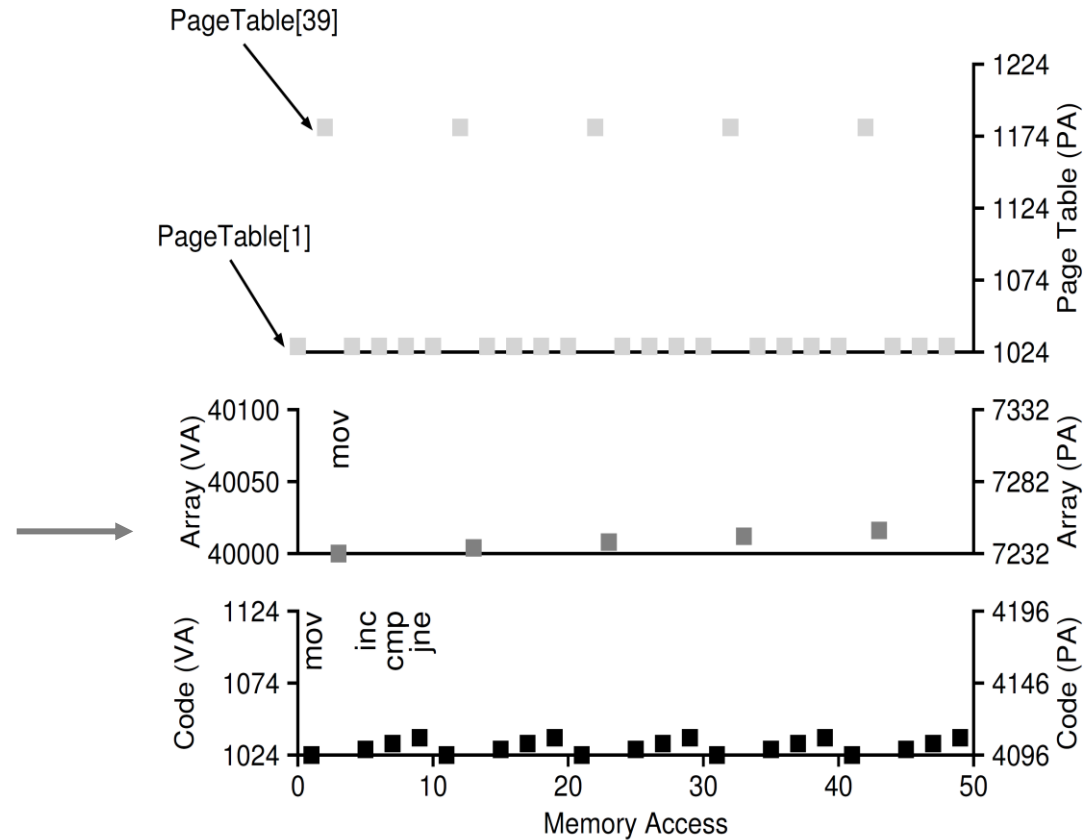
# Ex. Memory Trace

7

```
int array[1000];  
...  
for (i = 0; i < 1000; i++)  
    array[i] = 0;
```

↓

```
1024 movl $0x0, (%edi,%eax,4)  
1028 incl %eax  
1032 cmpl $0x03e8,%eax  
1036 jne 0x1024
```



Paging

ITP 30002  
Operating System

2021-04-14

# Translation-lookaside Buffer (TLB)

8

- A TLB is a part of MMU working as a cache of a page table
  - upon a memory access request, the computer architecture first checks the TLB to see if it has page translation information for the corresponding VPN
  - address-translation cache
- Steps for translating a virtual address with TLB
  - extract the VPN from a virtual address
  - check if the entry for the VPN is found in the TLB
    - if there exists, get the PFN from the TLB (i.e., TLB hit)
    - otherwise (i.e., TLB miss)
      - reference the page table to get the PFN
      - update the TLB with the VPN and the PFN
      - repeat from the beginning

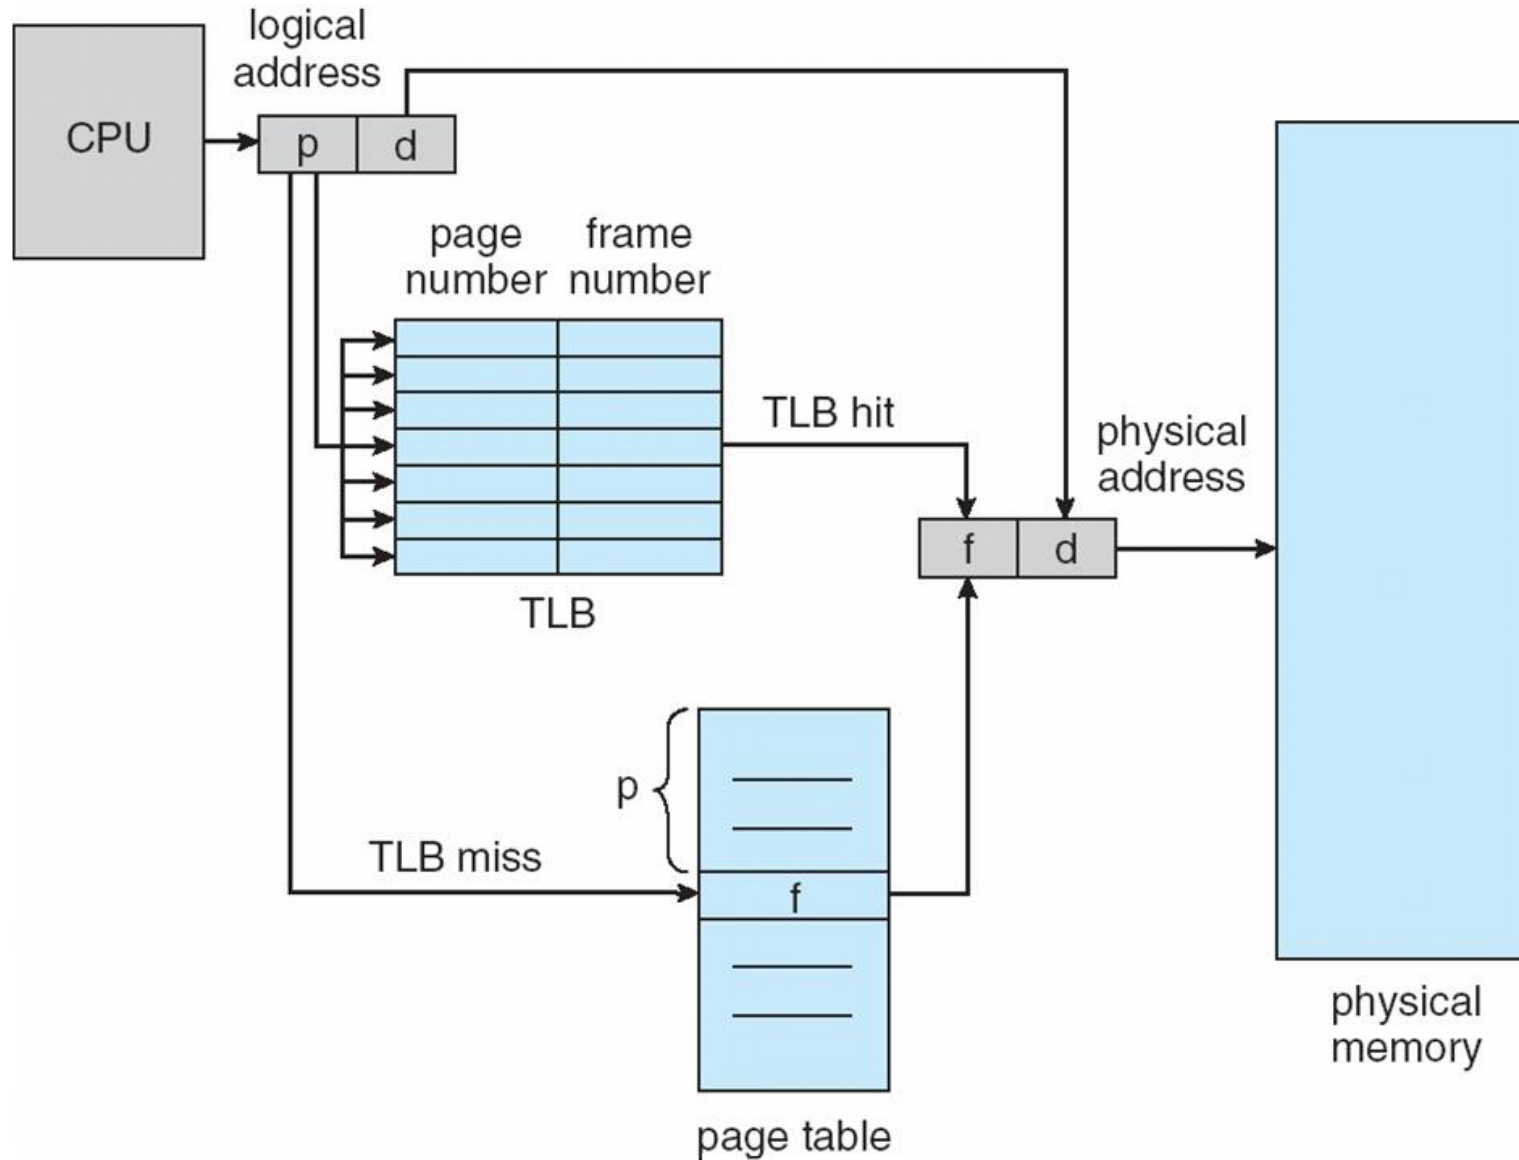
Paging

ITP 30002  
Operating System  
2021-04-14



# Paging Hardware with TLB

9



Paging

ITP 30002  
Operating System

2021-04-14

# Example

10

```
int sum = 0;
for (i = 0; i < 10; i++) {
    sum += a[i];
}
```

read 100 // a[0]

TLB miss

Update TLB

read 101 // a[1]

TLB hit

read 102 // a[2]

TLB hit

read 103 // a[3]

TLB miss

Update TLB


	Offset				
	00	04	08	12	16
VPN = 00					
VPN = 01					
VPN = 02					
VPN = 03					
VPN = 04					
VPN = 05					
VPN = 06		a[0]	a[1]	a[2]	
VPN = 07	a[3]	a[4]	a[5]	a[6]	
VPN = 08	a[7]	a[8]	a[9]		
VPN = 09					
VPN = 10					
VPN = 11					
VPN = 12					
VPN = 13					
VPN = 14					
VPN = 15					

Paging

ITP 30002  
Operating System

2021-04-14

# Memory Reference Pattern



11

Paging

ITP 30002  
Operating System

2021-04-14

# Locality

12

- A TLB can save memory accesses because memory accesses of an application program tend to have **temporal locality** and **spatial locality**
  - **temporal locality**: if a program accessed a memory at address  $x$ , it will likely access  $x$  again in near future.
  - **spatial locality**: if a program accesses a memory at address  $x$ , it will likely access  $x + d$  in near future (where  $d$  is a small number)
- Trade off between TLB size and TLB access speed
  - The bigger the size of a TLB is, the higher the hit ratio is
  - The bigger the size of a TLB is, the longer the look-up time takes

Paging

ITP 30002  
Operating System

2021-04-14