

ITP 30002 Operating System

Paging

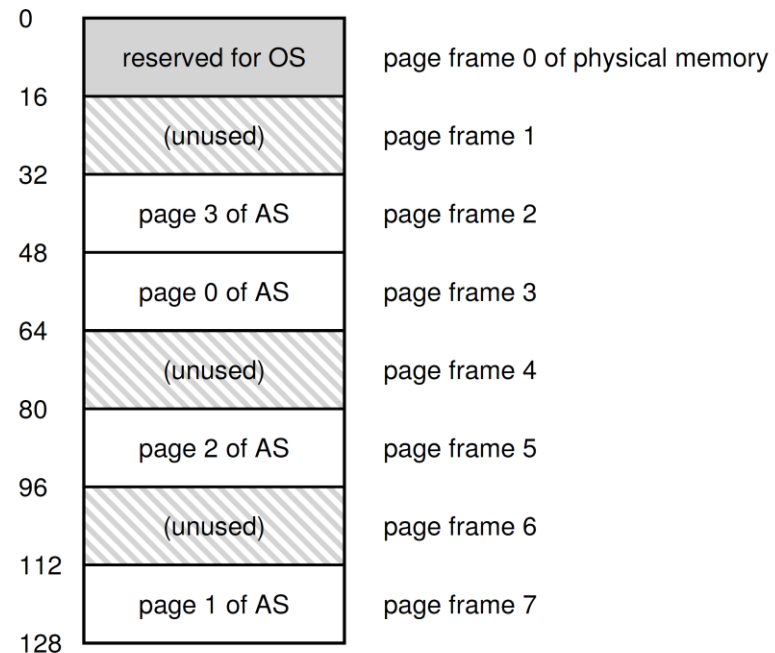
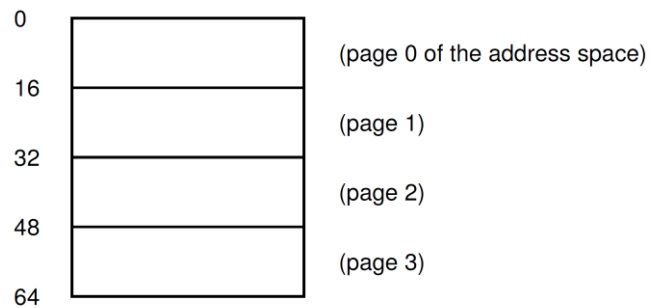
OSTEP Chapters 18 and 19

Shin Hong

Paging

2

- allocate memory to a process as a set of small fixed-size pieces
 - divide an address space into **pages** and a physical memory space into **frames** such that the sizes of a page and a frame are the same
 - can resolve both internal and external fragmentation problems (if the maintenance overhead is manageable)
- ex. 16-byte pages/frames



Paging

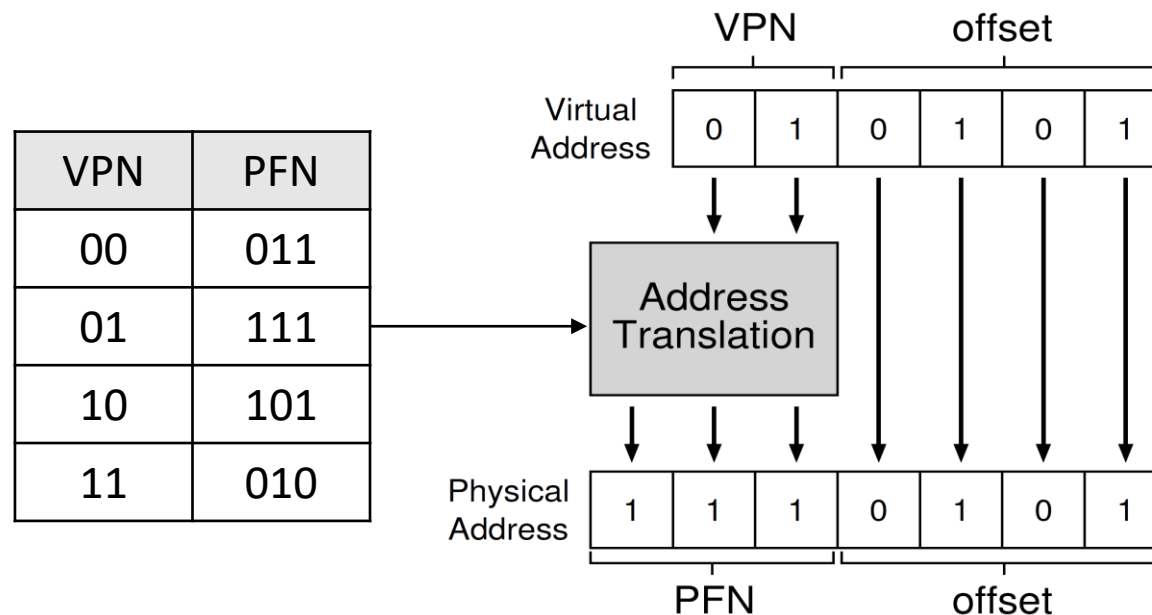
ITP 30002
Operating System

2021-04-17

Page Table and Address Translation

3

- page table records which frame a page of a process is assigned to
 - structured as per-process page table, or inverted page table
- for translation, a virtual address is splitted it into a VPN and an offset, and then the VPN is replaced with the corresponding PFN
 - ex. a 64-byte address space in a 128-byte physical memory



Paging

ITP 30002
Operating System

2021-04-17

Page Table Size

4

- page tables take a large amount of memory
 - example
 - 4 KB page/frame in a 32-bit address space and a 4 GB memory
 - 20-bits VPN, thus 2^{20} entries in a page table
 - 4 bytes for each entry, thus, 4 MB ($=2^{22}$ bytes) for a page table
 - 400 MB for 100 processes
- a page table should be resided in a main memory since MMU cannot accommodate a full page table
 - every memory access incurs at least one extra memory access for address translation

Paging

ITP 30002
Operating System


2021-04-17

Per-process Page Table Entry

5

- PFN
- valid bit: whether or not the page is in use
- present bit: whether the page is in physical memory or on disk
- permission bits: read-only, read-write, supermode-only, etc.
- dirty bit: whether the page was modified since it was brought into memory
- reference bit: whether the page has been recently accessed

- Example of 32-bit x86



Paging

ITP 30002
Operating System


2021-04-17

Accessing Memory with Paging

6

```
1 // Extract the VPN from the virtual address
2 VPN = (VirtualAddress & VPN_MASK) >> SHIFT
3
4 // Form the address of the page-table entry (PTE)
5 PTEAddr = PTBR + (VPN * sizeof(PTE))
6
7 // Fetch the PTE
8 PTE = AccessMemory(PTEAddr)
9
10 // Check if process can access the page
11 if (PTE.Valid == False)
12     RaiseException(SEGMENTATION_FAULT)
13 else if (CanAccess(PTE.ProtectBits) == False)
14     RaiseException(PROTECTION_FAULT)
15 else
16     // Access is OK: form physical address and fetch it
17     offset = VirtualAddress & OFFSET_MASK
18     PhysAddr = (PTE.PFN << PFN_SHIFT) | offset
19     Register = AccessMemory(PhysAddr)
```

VPN	PFN
00	011
01	111
10	101
11	010



Paging

ITP 30002
Operating System

2021-04-17


Ex. Memory Trace

7

```
int array[1000];  
...  
for (i = 0; i < 1000; i++)  
    array[i] = 0;
```

↓

```
1024 movl $0x0, (%edi,%eax,4)  
1028 incl %eax  
1032 cmpl $0x03e8,%eax  
1036 jne 0x1024
```



2
g System

Translation-lookaside Buffer (TLB)

8

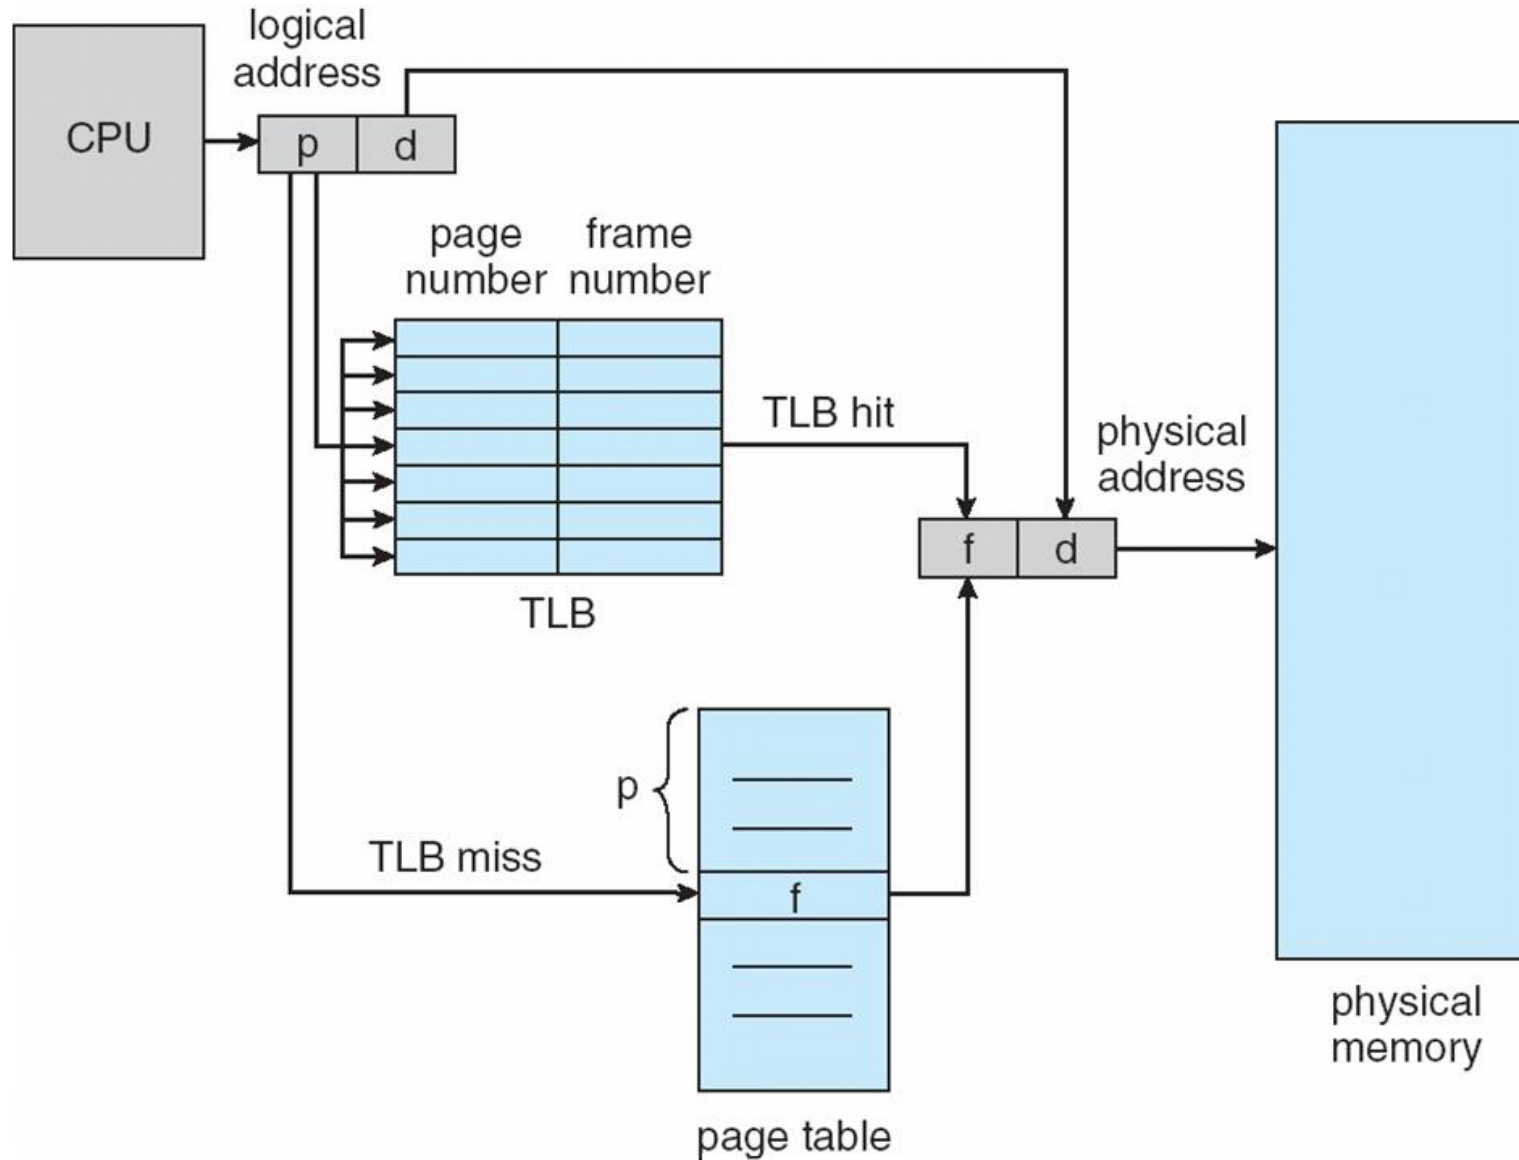
- A TLB is a part of MMU working as a cache of a page table
 - upon a memory access request, the computer architecture first checks the TLB to see if it has page translation information for the corresponding VPN
 - address-translation cache
- Steps for translating a virtual address with TLB
 - extract the VPN from a virtual address
 - check if the entry for the VPN is found in the TLB
 - if there exists, get the PFN from the TLB (i.e., TLB hit)
 - otherwise (i.e., TLB miss)
 - reference the page table to get the PFN
 - update the TLB with the VPN and the PFN
 - repeat from the beginning

Paging

ITP 30002
Operating System
2021-04-17

Paging Hardware with TLB

9



Paging

ITP 30002
Operating System

2021-04-17

Example

10

```
int sum = 0;
for (i = 0; i < 10; i++) {
    sum += a[i];
}
```

read 100 // a[0]

TLB miss

Update TLB

read 101 // a[1]

TLB hit

read 102 // a[2]

TLB hit

read 103 // a[3]

TLB miss

Update TLB


	Offset				
	00	04	08	12	16
VPN = 00					
VPN = 01					
VPN = 02					
VPN = 03					
VPN = 04					
VPN = 05					
VPN = 06		a[0]	a[1]	a[2]	
VPN = 07	a[3]	a[4]	a[5]	a[6]	
VPN = 08	a[7]	a[8]	a[9]		
VPN = 09					
VPN = 10					
VPN = 11					
VPN = 12					
VPN = 13					
VPN = 14					
VPN = 15					

Paging

ITP 30002
Operating System

2021-04-17

Memory Reference Pattern



11

Paging

ITP 30002
Operating System

2021-04-17

Locality

12

- A TLB can save memory accesses because memory accesses of an application program tend to have **temporal locality** and **spatial locality**
 - **temporal locality**: if a program accessed a memory at address x , it will likely access x again in near future.
 - **spatial locality**: if a program accesses a memory at address x , it will likely access $x + d$ in near future (where d is a small number)
- Trade off between TLB size and TLB access speed
 - The bigger the size of a TLB is, the higher the hit ratio is
 - The bigger the size of a TLB is, the longer the look-up time takes

Paging

ITP 30002
Operating System

2021-04-17

Translation-lookaside Buffer (TLB)

13

- A TLB is a part of MMU working as a cache of a page table
 - upon a memory access request, the computer architecture first checks the TLB to see if it has page translation information for the corresponding VPN
 - address-translation cache
- Steps for translating a virtual address with TLB
 - extract the VPN from a virtual address
 - check if the entry for the VPN is found in the TLB
 - if there exists, get the PFN from the TLB (i.e., TLB hit)
 - otherwise (i.e., TLB miss)
 - reference the page table to get the PFN
 - update the TLB with the VPN and the PFN
 - repeat from the beginning

Paging

ITP 30002
Operating System
2021-04-17

Handling TLB Miss

14

- Hardware-managed TLB
 - there must be a special register to point to the page table
 - the structure of a page table must be fixed
- Software-managed TLB
 - the architecture raises an exception in a middle of an instruction execution if a TLB miss occurs
 - the OS finds proper information from page tables and updates TLB
 - the architecture retries the instruction

Paging

ITP 30002
Operating System

2021-04-17

TLB Entry

15

- A TLB of an architecture typically has 32, 64 or 128 entries
 - parallel search
- A TLB entry has the following attributes
 - VPN
 - PFN
 - valid bit
 - protection bit: read-only or read-write
 - dirty bit

Paging

ITP 30002
Operating System

2021-04-17

Managing TLB at Context Switching

16

- The existing TLB entries must not be used after a context switching
- Approaches
 1. Flush all entries
 - flushing can be made by turning off the valid bit of every entry
 - a series of TLB misses will happen right after a context switch
 2. Have a process identifier in a TLB entry
 - address space identifier (ASID)
 - multiple processes can readily share a TLB

Paging

ITP 30002
Operating System

2021-04-17

TLB Entry Replacement Policies

17

- Which entry to evict when the TLB is full and a new entry must be installed?
 - decision should be made to minimize a TLB miss rate
 - a typical case of the cache replacement problem
- Approaches
 1. evict the least-recently-used (LRU) one
 2. evict a random entry

Paging

ITP 30002
Operating System

2021-04-17