

ITP 30002 Operating System

# Paging: Smaller Tables

OSTEP Chapter 20

Shin Hong

# Make Page Table Smaller

2

- Array-based page tables take too much memory resource even though most of page entries are invalid
  - e.g., for a 32-bit address space with 4 KB pages, a per-process page table takes 4 MB
- Approaches
  1. use bigger pages
  2. per-segment page tables
  3. multi-level page tables
  4. inverted page table

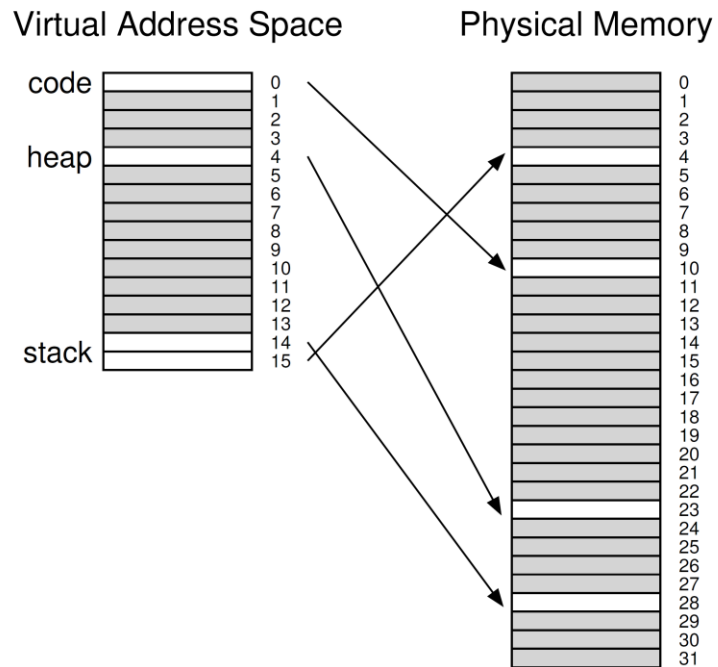
Paging: Smaller  
Page Tables

ITP 30002  
Operating System

2021-05-14

# Per-segment Page Table

3



| PFN | valid | prot | present | dirty |
|-----|-------|------|---------|-------|
| 10  | 1     | r-x  | 1       | 0     |
| -   | 0     | —    | -       | -     |
| -   | 0     | —    | -       | -     |
| -   | 0     | —    | -       | -     |
| 23  | 1     | rw-  | 1       | 1     |
| -   | 0     | —    | -       | -     |
| -   | 0     | —    | -       | -     |
| -   | 0     | —    | -       | -     |
| -   | 0     | —    | -       | -     |
| -   | 0     | —    | -       | -     |
| -   | 0     | —    | -       | -     |
| -   | 0     | —    | -       | -     |
| -   | 0     | —    | -       | -     |
| -   | 0     | —    | -       | -     |
| 28  | 1     | rw-  | 1       | 1     |
| 4   | 1     | rw-  | 1       | 1     |

- Based on the observation that only first few pages of each segment is used in many programs
- Have three page tables for three segments, and allocate a variable-length memory region for holding a page table
  - the base register points to the beginning of a page table, and the bounds register represents the number of allocated pages in the segment

Paging: Smaller  
Page Tables

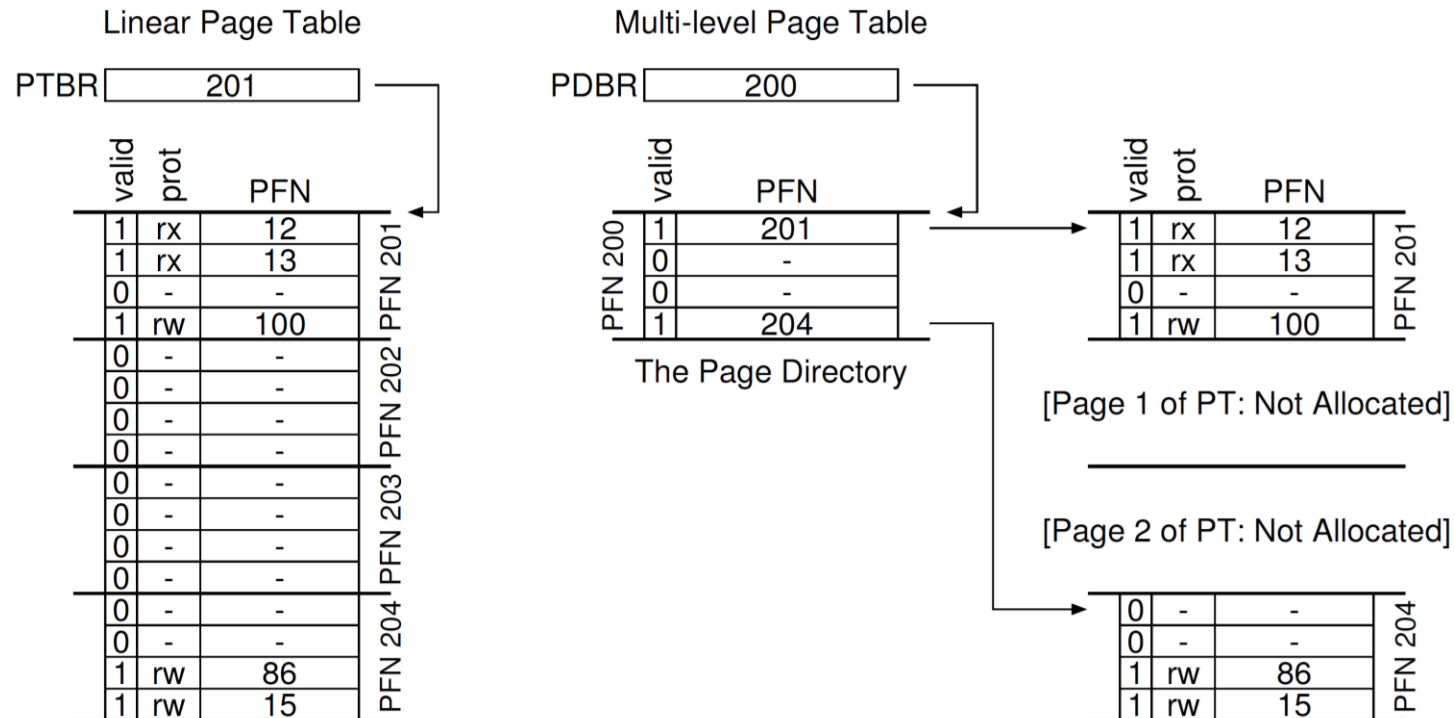
ITP 30002  
Operating System  
2021-05-14



# Multi-level Page Tables

5

- Divide a page table into page-size units
- Allocate a page to each page table piece, but do not allocate a page if the corresponding unit has no valid entry
- Have a **page directory** as an index of the allocated pages for a page table
- E.g.



Paging: Smaller  
Page Tables

ITP 30002  
Operating System

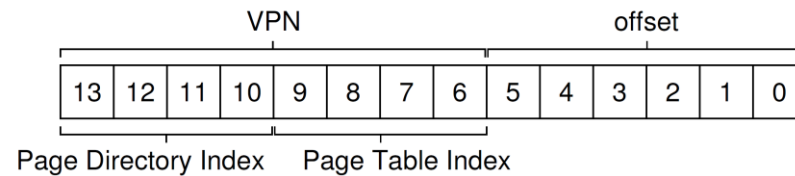
2021-05-14

# Example

6

- Address size of 16 KB ( $2^{14}$ ) with 64-byte pages
  - 256 VPNs
  - a page table takes 1 KB (16 pages) if a PTE takes 4 bytes
- Suppose that VPNs only 0, 1, 4, 5, 254 and 255 are used, and the rest are unused

|           |                  |
|-----------|------------------|
| 0000 0000 | code             |
| 0000 0001 | code             |
| 0000 0010 | (free)           |
| 0000 0011 | (free)           |
| 0000 0100 | heap             |
| 0000 0101 | heap             |
| 0000 0110 | (free)           |
| 0000 0111 | (free)           |
| .....     | ... all free ... |
| 1111 1100 | (free)           |
| 1111 1101 | (free)           |
| 1111 1110 | stack            |
| 1111 1111 | stack            |



$$\text{PDEAddr} = \text{PageDirBase} + \text{PDIndex} * \text{sizeof(PDE)}$$

$$\text{PTEAddr} = (\text{PDEAddr} \rightarrow \text{PFN} \ll \text{SHIFT}) + \text{PTIndex} * \text{sizeof(PTE)}$$

| Page Directory<br>PFN | valid? |
|-----------------------|--------|
| 100                   | 1      |
| —                     | 0      |
| —                     | 0      |
| —                     | 0      |
| —                     | 0      |
| —                     | 0      |
| —                     | 0      |
| —                     | 0      |
| —                     | 0      |
| —                     | 0      |
| —                     | 0      |
| —                     | 0      |
| —                     | 0      |
| —                     | 0      |
| —                     | 0      |
| —                     | 0      |
| 101                   | 1      |

| Page of PT (@PFN:100)<br>PFN | valid | prot |
|------------------------------|-------|------|
| 10                           | 1     | r-x  |
| 23                           | 1     | r-x  |
| —                            | 0     | —    |
| —                            | 0     | —    |
| 80                           | 1     | rw-  |
| 59                           | 1     | rw-  |
| —                            | 0     | —    |
| —                            | 0     | —    |
| —                            | 0     | —    |
| —                            | 0     | —    |
| —                            | 0     | —    |
| —                            | 0     | —    |
| —                            | 0     | —    |
| —                            | 0     | —    |
| —                            | 0     | —    |
| —                            | 0     | —    |
| —                            | 0     | —    |
| —                            | 0     | —    |

| Page of PT (@PFN:101)<br>PFN | valid | prot |
|------------------------------|-------|------|
| —                            | 0     | —    |
| —                            | 0     | —    |
| —                            | 0     | —    |
| —                            | 0     | —    |
| —                            | 0     | —    |
| —                            | 0     | —    |
| —                            | 0     | —    |
| —                            | 0     | —    |
| —                            | 0     | —    |
| —                            | 0     | —    |
| —                            | 0     | —    |
| —                            | 0     | —    |
| —                            | 0     | —    |
| —                            | 0     | —    |
| 55                           | 1     | rw-  |
| 45                           | 1     | rw-  |

# Two-level Page Table Control Flow

7

```
1  VPN = (VirtualAddress & VPN_MASK) >> SHIFT
2  (Success, TlbEntry) = TLB_Lookup(VPN)
3  if (Success == True)    // TLB Hit
4      if (CanAccess(TlbEntry.ProtectBits) == True)
5          Offset    = VirtualAddress & OFFSET_MASK
6          PhysAddr = (TlbEntry.PFN << SHIFT) | Offset
7          Register = AccessMemory (PhysAddr)
8      else
9          RaiseException(PROTECTION_FAULT)
10 else    // TLB Miss
11     // first, get page directory entry
12     PDIndex = (VPN & PD_MASK) >> PD_SHIFT
13     PDEAddr = PDBR + (PDIndex * sizeof(PDE))
14     PDE      = AccessMemory (PDEAddr)
15     if (PDE.Valid == False)
16         RaiseException(SEGMENTATION_FAULT)
17     else
18         // PDE is valid: now fetch PTE from page table
19         PTIndex = (VPN & PT_MASK) >> PT_SHIFT
20         PTEAddr = (PDE.PFN << SHIFT) + (PTIndex * sizeof(PTE))
21         PTE      = AccessMemory (PTEAddr)
22         if (PTE.Valid == False)
23             RaiseException(SEGMENTATION_FAULT)
24         else if (CanAccess(PTE.ProtectBits) == False)
25             RaiseException(PROTECTION_FAULT)
26         else
27             TLB_Insert(VPN, PTE.PFN, PTE.ProtectBits)
28             RetryInstruction()
```

Paging: Smaller  
Page Tables

ITP 30002  
Operating System

2021-05-14

# Inverted Page Table

8

- Keep a single page table that has an entry for each frame
  - each frame is mapped to a VPN of a process
  - only one page table would be enough for a single system
- Searching the entry for a VPN of a process consumes much more time than array-based page tables
  - linear search, hashing, etc.

Paging: Smaller  
Page Tables

ITP 30002  
Operating System

2021-05-14