

ITP 30002 Operating System

Files and Directories

Chapters 39

Shin Hong

Two Abstractions for Persistent Storage

2

- A **file** is a linear array of bytes, each of which can be read or write
 - each file has a low-level identifier often called **inode number**
 - the structure inside a file is managed by the application program, but the file's persistency is managed by the OS
- A **directory** is a list of names of the contained files and directories
 - a name is a pair of user-readable one and low-level one
 - a directory is named by an inode number
- The directory hierarchy is formed as a directory exists inside another one
 - the hierarchy starts at a root directory (i.e., /)
 - the **file location** is defined by the concatenation of all directory names from the root down to the immediate container and the file name

Files and
Directories

ITP 30002
Operating System

2021-06-14

File System Interface – File Access (1)

3

- Open a File (i.e., `open()`)
 - create a new file, open an existing one, truncate a file to a size of zero
 - return a file descriptor which is a non-negative integer private per process, works as a pointer to a file object
- Read or write a file sequentially
 - example:

```
prompt> strace cat foo
...
open("foo", O_RDONLY|O_LARGEFILE)      = 3
read(3, "hello\n", 4096)                = 6
write(1, "hello\n", 6)                  = 6
hello
read(3, "", 4096)                       = 0
close(3)                                = 0
...
prompt>
```

Files and
Directories


ITP 30002
Operating System

2021-06-14

File Access (2)

4

- Change the current offset of a file (i.e., `lseek()`)
 - once a file is opened, the OS tracks the current offset to determine which offset will be read or written in the next
 - `lseek` is to explicitly move an file offset to a specific point within a file to read or write to random offsets
- Share file table entries
 - a process has a unique entry in the open file table as it opens and reads/writes a file
 - a child processes created by `fork()` shares the same open file entry with its parent process
 - `dup()` allows a process to create a new file descriptor that refers to the same underlying open file as an existing file descriptor



File Access (3)

5

- Flush (i.e., `fsync()`)
 - enforce the file system to write the buffered data of the given file descriptor to the persistent storage
- Get information about files (i.e., `stat()`, `fstat()`)
 - the metadata for a file contains the size of a file, the low-level name, ownership information, when the file was accessed or modified, etc.
 - the metadata of a file is stored in an inode structure.

```
struct stat {
    dev_t      st_dev;      // ID of device containing file
    ino_t      st_ino;      // inode number
    mode_t     st_mode;     // protection
    nlink_t    st_nlink;    // number of hard links
    uid_t      st_uid;      // user ID of owner
    gid_t      st_gid;      // group ID of owner
    dev_t      st_rdev;     // device ID (if special file)
    off_t      st_size;     // total size, in bytes
    blksize_t  st_blksize;  // blocksize for filesystem I/O
    blkcnt_t   st_blocks;   // number of blocks allocated
    time_t     st_atime;    // time of last access
    time_t     st_mtime;    // time of last modification
    time_t     st_ctime;    // time of last status change
};
```

Files and
Directories

ITP 30002
Operating System

2021-06-14

Directory

6

- Create an empty directory (i.e., `mkdir()`)
 - an empty directory contains two entries (i.e., `.` and `..`) by default.
 - each directory entry, which works as a link to a file/directory, is structured as follows:

```
struct dirent {  
    char            d_name[256]; // filename  
    ino_t           d_ino;       // inode number  
    off_t           d_off;       // offset to the next dirent  
    unsigned short  d_reclen;    // length of this record  
    unsigned char    d_type;     // type of file  
};
```

- Create a new directory entry from an existing one (i.e., `ln()`)
 - create a new file name (i.e., hard link) to refer the inode of the same existing file
 - Example.

```
prompt> ls -i file file2  
67158084 file  
67158084 file2
```
 - a soft link is a special file that points to another file name
- Delete a file from a directory
 - unlink the corresponding link from a directory (i.e., the reference count is decreased)
 - the file system will eventually reclaim if an inode has zero count

Files and
Directories

ITP 30002
Operating System

2021-06-14

Permission Bits and Access Control

7

- Unlike process and virtual memory, file systems are basically shared among multiple users, thus they require more ways to control permissions
- Permission bit
 - the owner of a file can update the permission
 - (owner, group, others) X (read, write, execute)
 - for a directory, an execute bit allows the user to change directory into the one
 - example:

```
prompt> ls -l foo.txt  
-rw-r--r--  1 remzi wheel  0 Aug 24 16:29 foo.txt
```

- Access control list
 - specify what kinds of rights a specific user has on specific files

Files and
Directories

ITP 30002
Operating System

2021-06-14

Making and Mounting A File System

8

- Make a file system
 - give a file system creating tool (e.g., mkfs), as input, a device and a file system type.
- Mount a file system
 - take an existing directory as a target **mount point** and paste a new file system onto the directory free at that point
 - example

```
prompt> mount -t ext3 /dev/sda1 /home/users  
prompt> ls /home/users/  
a b
```

Files and
Directories

ITP 30002
Operating System

2021-06-14