

ITP 30002 Operating System

# Files and Directories

Chapters 39

Shin Hong

# Two Abstractions for Persistent Storage

2

- A **file** is a linear array of bytes, each of which can be read or write
  - each file has a low-level identifier often called **inode number**
  - the structure inside a file is managed by the application program, but the file's persistency is managed by the OS
- A **directory** is a list of names of the contained files and directories
  - a name is a pair of user-readable one and low-level one
  - a directory is named by an inode number
- The directory hierarchy is formed as a directory exists inside another one
  - the hierarchy starts at a root directory (i.e., /)
  - the **file location** is defined by the concatenation of all directory names from the root down to the immediate container and the file name

Files and  
Directories

ITP 30002  
Operating System

2023-06-06

# File System Interface – File Access (1)

3

- Open a File (i.e., `open()`)
  - create a new file, open an existing one, truncate a file to a size of zero
  - return a file descriptor which is a non-negative integer private per process, works as a pointer to a file object
- Read or write a file sequentially
  - example:

```
prompt> strace cat foo
...
open("foo", O_RDONLY|O_LARGEFILE)      = 3
read(3, "hello\n", 4096)                = 6
write(1, "hello\n", 6)                  = 6
hello
read(3, "", 4096)                       = 0
close(3)                                = 0
...
prompt>
```

Files and  
Directories


ITP 30002  
Operating System

2023-06-06

# File Access (2)

4

- Change the current offset of a file (i.e., `lseek()`)
  - once a file is opened, the OS tracks the current offset to determine which offset will be read or written in the next
  - `lseek` is to explicitly move a file offset to a specific point within a file to read or write to random offsets
- Share file table entries
  - a process has a unique entry in the open file table as it opens and reads/writes a file
  - a child processes created by `fork()` shares the same open file entry with its parent process
  - `dup()` allows a process to create a new file descriptor that refers to the same underlying open file as an existing file descriptor



# File Access (3)

5

- Flush (i.e., `fsync()`)
  - enforce the file system to write the buffered data of the given file descriptor to the persistent storage
- Get information about files (i.e., `stat()`, `fstat()`)
  - the metadata for a file contains the size of a file, the low-level name, ownership information, when the file was accessed or modified, etc.
  - the metadata of a file is stored in an inode structure.

```
struct stat {  
    dev_t    st_dev;        // ID of device containing file  
    ino_t    st_ino;        // inode number  
    mode_t   st_mode;       // protection  
    nlink_t  st_nlink;      // number of hard links  
    uid_t    st_uid;        // user ID of owner  
    gid_t    st_gid;        // group ID of owner  
    dev_t    st_rdev;       // device ID (if special file)  
    off_t    st_size;       // total size, in bytes  
    blksize_t st_blksize;   // blocksize for filesystem I/O  
    blkcnt_t st_blocks;     // number of blocks allocated  
    time_t   st_atime;      // time of last access  
    time_t   st_mtime;      // time of last modification  
    time_t   st_ctime;      // time of last status change  
};
```

Files and  
Directories

ITP 30002  
Operating System

2023-06-06

# Directory

6

- Create an empty directory (i.e., `mkdir()`)
  - an empty directory contains two entries (i.e., `.` and `..`) by default.
  - each directory entry, which works as a link to a file/directory, is structured as follows:

```
struct dirent {  
    char            d_name[256]; // filename  
    ino_t           d_ino;       // inode number  
    off_t           d_off;       // offset to the next dirent  
    unsigned short  d_reclen;    // length of this record  
    unsigned char    d_type;     // type of file  
};
```

- Create a new directory entry from an existing one (i.e., `ln()`)
  - create a new file name (i.e., hard link) to refer the inode of the same existing file
    - Example. 

```
prompt> ls -i file file2  
67158084 file  
67158084 file2
```
  - a soft link is a special file that points to another file name
- Delete a file from a directory
  - unlink the corresponding link from a directory (i.e., the reference count is decreased)
  - the file system will eventually reclaim if an inode has zero count

Files and  
Directories

ITP 30002  
Operating System  
2023-06-06

# Permission Bits and Access Control

7

- Unlike process and virtual memory, file systems are basically shared among multiple users, thus they require more ways to control permissions
- Permission bit
  - the owner of a file can update the permission
  - (owner, group, others) X (read, write, execute)
  - for a directory, an execute bit allows the user to change directory into the one
  - example:

```
prompt> ls -l foo.txt
-rw-r--r--  1 remzi wheel  0 Aug 24 16:29 foo.txt
```

- Access control list
  - specify what kinds of rights a specific user has on specific files

Files and  
Directories

ITP 30002  
Operating System

2023-06-06

# Making and Mounting A File System

8

- Make a file system
  - give a file system creating tool (e.g., mkfs), as input, a device and a file system type.
- Mount a file system
  - take an existing directory as a target **mount point** and paste a new file system onto the directory free at that point
  - example

```
prompt> mount -t ext3 /dev/sda1 /home/users  
prompt> ls /home/users/  
a b
```

Files and  
Directories

ITP 30002  
Operating System

2023-06-06