

Deadlock and its Detection

Shin Hong

Deadlock Bugs Frequently Occur in Real World

Application	What it does	Non-Deadlock	Deadlock
MySQL	Database Server	14	9
Apache	Web Server	13	4
Mozilla	Web Browser	41	16
OpenOffice	Office Suite	6	2
Total		74	31

- In a survey on 105 real-world concurrency bugs in open-source applications, **31 out of 105 bugs are deadlock bugs** [Lu *et al.*, ASPLOS 08]


Deadlock

- A deadlock occurs when each of a set of threads is blocked, waiting for another thread in the set to satisfy **certain condition**
 - **release shared resource**
 - **raise event**

Resource Deadlock in Concurrent Programs

- ABBA deadlock

Thread1 () {	Thread2 () {
1: lock (X)	11: lock (Y)
2: x = ... ;	12: y = ... ;
3: lock (Y)	13: lock (X)
4: y = ... ;	14: x = ... ;
5: unlock (Y)	15: unlock (X)
6: unlock (X)	16: unlock (Y)
}	}

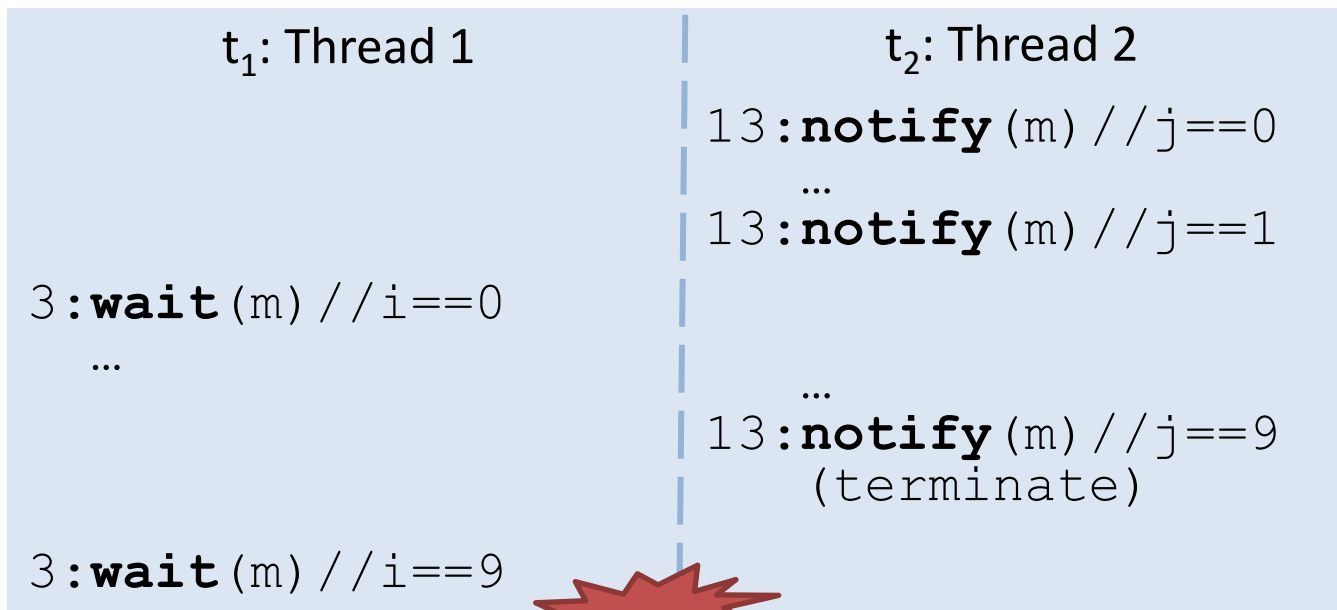


Communication Deadlock

- Lost notify

```
Thread1 () {  
1: ...  
2: for (i=0; i<10; i++) {  
3:   wait (m) ; }  
}
```

```
Thread2 () {  
11: ...  
12: for (j=0; j<10; j++) {  
13:   notify (m) ; }  
}
```



Finding Deadlock Bugs is Difficult

- A deadlock bug induces deadlock situations **only under certain thread schedules**
- Systems software creates a **massive number of locks** for fine-grained concurrency controls
- **Function caller-callee relation** complicates the reasoning about possible nested lockings

Bug Detection Approach

Resource deadlock

- Basic potential deadlock detection algorithm
- GoodLock algorithm

Communication deadlock

- CHECKMATE: a trace program model-checking technique for deadlock detection


Basic Potential Deadlock Detection

- Extend the cyclic deadlock monitoring algorithm
 - Cyclic deadlock monitoring algorithm (e.g. *LockDep*)
 - Monitor lock acquires and releases in runtime
 - Lock graph (N, E_N)
 - Create a node n_X when a thread acquires lock X
 - Create an edge (n_X, n_Y) when a thread acquires lock Y while holding lock X
 - Remove n_X , $(n_X, *)$ and $(*, n_X)$ when a thread releases X
- Report deadlock when the graph has any cycle

Cyclic Deadlock Detection Example (1/2)

```
Thread1 () {  
1: lock (X)  
2: a = ... ;  
3: lock (Y)  
4: b = ... ;  
5: unlock (Y)  
6: unlock (X)  
}  
  
Thread2 () {  
11: lock (Y)  
12: b = ... ;  
13: lock (X)  
14: a = ... ;  
15: unlock (X)  
16: unlock (Y)  
}
```

<u>t1: Thread 1</u>	<u>t2: Thread 2</u>
1: lock (X)	
2: a = ...	
3: lock (Y)	11: lock (Y)
	12: b = ...
	13: lock (X)



Deadlock detected!

Cyclic Deadlock Detection Example (2/2)

```
Thread1 () {      Thread2 () {
1: lock (X) ;      11: lock (Y) ;
2: a = ...         12: b = ...
3: lock (Y) ;      13: lock (X) ;
4: b = ...         14: a = ...
5: unlock (Y) ;    15: unlock (X) ;
6: unlock (X) ;    16: unlock (Y) ;
}                  }
```

<u>t1: Thread 1</u>	<u>t2: Thread 2</u>
1: lock (X)	
2: a = ...	
3: lock (Y)	
4: b = ...	
5: unlock (Y)	
6: unlock (X)	
	11: lock (Y)
	12: b = ...
	13: lock (X)
	14: a = ...
	15: unlock (X)
	16: unlock (Y)

No problem

Basic Deadlock Prediction Technique


- Potential cyclic deadlock detection algorithm [Harrow, SPIN 00]
 - Lock graph (N, E_N)
 - Create a node n_x when a thread acquires lock X
 - Create an edge (n_x, n_y) when a thread acquires lock Y while holding lock X
 - ~~• Remove $n_x, (n_x, *)$ and $(*, n_x)$ when a thread releases X~~
 - Report potential deadlocks if the resulted graph at the end of an execution has a cycle

[Harrow, SPIN 00] J. J. Harrow, Jr.: Runtime checking of multithreaded applications with Visual Threads, SPIN Workshop 2000

Potential Cyclic Deadlock Detection Example

```
Thread1 () {  
1: lock (X)  
2: a = ... ;  
3: lock (Y)  
4: b = ... ;  
5: unlock (Y)  
6: unlock (X)  
}  
  
Thread2 () {  
11: lock (Y)  
12: b = ... ;  
13: lock (X)  
14: a = ... ;  
15: unlock (X)  
16: unlock (Y)  
}
```

<u>t1:Thread 1</u>	<u>t2:Thread 2</u>
1:lock (X)	
2:a = ...	
3:lock (Y)	
4:b = ...	
5:unlock (Y)	
6:unlock (X)	11:lock (Y)
	12:b=...
	13:lock (X)
	...




Cycle → Potential deadlock

Basic Deadlock Prediction Technique

- The algorithm is commercialized as a SW tool VisualThreads (*HP*)
- Empirical results show that the algorithm is very effective to discover hidden deadlock bugs
- Challenge: generate many false positive

False Positive Example#1 – Single Thread Cycle

```
Thread1 () {  
1: lock (X) ;  
2: lock (Y) ;  
3: unlock (Y) ;  
4: unlock (X) ;  
  
5: lock (Y) ;  
6: lock (X) ;  
7: unlock (X) ;  
8: unlock (Y) ;}  
  
Thread2 () {  
11: lock (X) ;  
12: unlock (X) ;  
  
13: lock (Y) ;  
14: unlock (Y) ;}
```



The lock graph has a cycle, but no deadlock

A cycle that consists of edges created by one thread is a false positive

False Positive Example#2: Gate Lock

Thread1 () {

1: lock (X) ;

2: lock (Y) ;

3: lock (Z) ;

4: unlock (Z) ;

5: unlock (Y) ;

6: unlock (X) ; }

Thread2 () {

11: lock (X) ;

12: lock (Z) ;


13: lock (Y) ;

14: unlock (Y) ;

15: unlock (Z) ;

16: unlock (X) ; }

Gate lock
(guard lock)




Cycle, but no deadlock

False Positive Example#3: Thread Creation

```
main() {  
0:  start(f1);  
}  
  
f1() {  
1:  lock(X);  
2:  lock(Y);  
3:  unlock(Y);  
4:  unlock(X);  
5:  start(f2);  
}  
  
f2() {  
11: lock(Y);  
12: lock(X);  
13: unlock(X);  
14: unlock(Y);  
}
```

Thread segment#1

Thread segment#2



Cycle, but no deadlock

GoodLock Algorithm[Agarwal, IBM 10]

- Extend the lock graph in the basic potential deadlock detection algorithm to consider *thread*, *gate lock*, and *thread segment*
- A cycle is *valid* (i.e., true positive) when every pair of edges $(m_{11}, (s_{11}, t_1, G_1, s_{12}), m_{12})$, and $(m_{21}, (s_{21}, t_2, G_2, s_{22}), m_{22})$ in the cycle satisfies:
 - $t_1 \neq t_2$, and
 - $G_1 \cap G_2 = \emptyset$, and
 - $\neg(s_{12} < s_{21})$
 - The happens-before relation $<$ is the transitive closure of the relation R such that $(s_1, s_2) \in R$ if there exists the edge from s_1 to s_2 in the thread segment graph

[Agarwal, IBM 10] R. Agarwal et al., Detection of deadlock potential in multithreaded programs, IBM Journal of Research and Development, 54(5), 2010