

ITP30002 Operating System

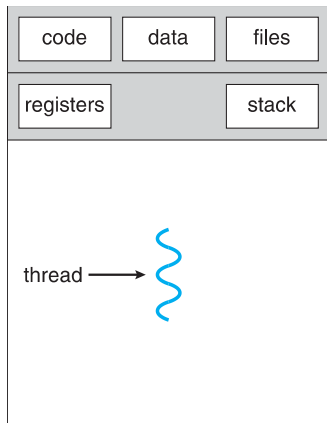
Multithreading

This lecture note is taken from the instructor's resource of Operating System Concept, 10/e and then partly edited/revised by Shin Hong.

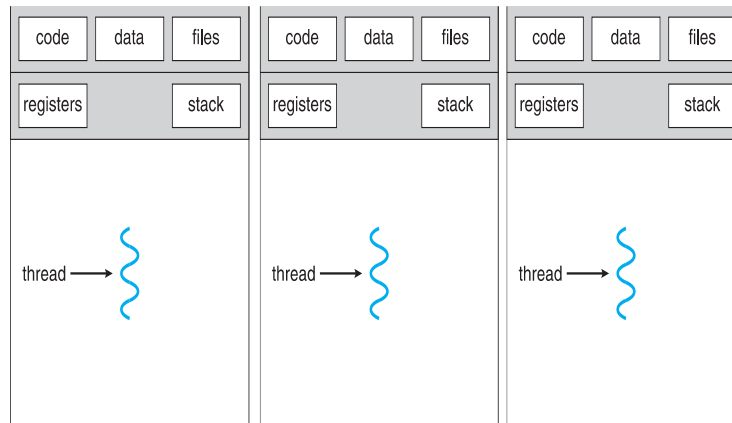
What is Multithreading?

- A thread is a single flow of control
- Threads run within an application
- Most modern applications are multithreaded to run multiple tasks within an application
 - E.g., a web browser does update display, fetch data, spell checking, and answer a network request at the same time within an application


Single and Multithreaded Processes



<Single threaded process>



<Multi-processing>




<Multi-threaded process>

Benefits

- Scalability: process can take advantage of multiprocessor architectures
- Responsiveness: may allow continued execution if part of process is blocked, especially important for user interfaces
- Resource sharing: threads share resources of process, easier than shared memory or message passings
- Efficiency: cheaper than process creation, thread switching lower overhead than context switching


Why Concurrent Programming




- *Parallelism* implies a system can perform more than one task simultaneously
- *Concurrency* supports more than one task making progress simultaneously
 - Single processor/core, scheduler providing concurrency
- Multicore or multiprocessor systems putting pressure on programmers, challenges include:
 - Dividing activities
 - Balance
 - Data splitting
 - Data dependency
 - Testing and debugging

Concurrency vs. Parallelism

Concurrent execution on single-core system:



Parallelism on a multi-core system:



Multicore Programming (Cont.)

- Types of parallelism
 - Data parallelism – distributes subsets of the same data across multiple cores, same operation on each
 - Task parallelism – distributing threads across cores, each thread performing unique operation
- As the number of threads grows in programs, so does architectural support for threading
 - CPUs have cores as well as *hardware threads*

Amdahl's Law


- Upper bound of performance gains by adding additional cores to an application that has both serial and parallel components
 - S is serial portion
 - N processing cores
- $$speedup \leq \frac{1}{S + \frac{(1-S)}{N}}$$
- That is, if application is 75% parallel / 25% serial, moving from 1 to 2 cores results in speedup of 1.6 times
 - As N approaches infinity, speedup approaches $1 / S$
- Serial portion of an application has disproportionate effect on performance gained by adding additional cores

User Threads and Kernel Threads

- User threads: management done by user-level threads library
 - POSIX Pthreads
 - Windows threads
 - Java threads
- Kernel threads: supported by the Kernel
 - Windows
 - Solaris
 - Linux
 - Tru64 UNIX
 - Mac OS X


Many-to-One

- Many user-level threads mapped to single kernel thread
- One thread blocking causes all to block
- Multiple threads may not run in parallel on muticore system because only one may be in kernel at a time
- Few systems currently use this model
 - Solaris Green Threads
 - GNU Portable Threads




One-to-One

- Each user-level thread maps to a kernel thread
- Creating a user-level thread creates a kernel thread
- More concurrency than many-to-one
- Number of threads per process sometimes restricted due to overhead
 - Windows
 - Linux
 - Solaris 9 and later




Many-to-Many Model

- Allows many user level threads to be mapped to many kernel threads
- Allows the operating system to create a sufficient number of kernel threads
 - Solaris prior to version 9
 - Windows with the *ThreadFiber* package



Two-level Model

- Similar to M:M, except that it allows a user thread to be **bound** to kernel thread
- Examples
 - IRIX
 - HP-UX
 - Tru64 UNIX
 - Solaris 8 and earlier



Thread Libraries

- **Thread library** provides programmer with API for creating and managing threads
- Two primary ways of implementing
 - Library entirely in user space
 - Kernel-level library supported by the OS

Pthread (POSIX Thread)

- A POSIX standard (IEEE 1003.1c) API for thread creation and synchronization
 - Specification, not implementation;
 - API specifies behavior of the thread library, implementation is up to development of the library
 - May be provided either as user-level or kernel-level
- Common in UNIX operating systems

Implicit Threading

- Growing in popularity as the number of threads increases, ensuring program correctness is more difficult with explicit threads
- Creation and management of threads done by compilers and run-time libraries rather than programmers
 - OpenMP
 - Thread Pools
 - Grand Central Dispatch
 - Microsoft Threading Building Blocks (TBB)
 - `java.util.concurrent` package

- A set of compiler directives and an API for C, C++, FORTRAN
- Provides support for parallel programming in shared-memory environments
- Identifies **parallel regions** – blocks of code that can run in parallel
 - run for loop in parallel
 - create as many threads as there are cores
 - ex.

```
#pragma omp parallel for for(i=0;i<N;i++) {  
    c[i] = a[i] + b[i];  
}
```

Thread Pools

- Create a number of threads in a pool where they await work
- Advantages
 - Slightly faster to service a request with an existing thread than create a new thread
 - Allows the number of threads in the application(s) to be bound to the size of the pool
 - Separating parallel tasks from threading mechanisms

Threading Design Issue: Fork & Exec

- Does **fork ()** duplicate only the calling thread or all threads?
 - Some UNIXes have two versions of fork
- **exec ()** usually works as to replace the running process including all threads

Threading Design Issue: Signal

- Where should a signal be delivered for multi-threaded?
 - Deliver the signal to the thread to which the signal applies
 - Deliver the signal to every thread in the process
 - Deliver the signal to certain threads in the process
 - Assign a specific thread to receive all signals for the process

Threading Design Issue: Thread Cancellation

- Terminating a thread before it has finished
- Two general approaches
 - Asynchronous cancellation terminates the target thread immediately
 - Deferred cancellation allows the target thread to periodically check if it should be cancelled
 - a thread terminates next time when it reaches at a predefined cancellation point

Threading Design Issue: Thread-Local Storage

- Thread-local storage (TLS) allows each thread to have its own copy of data
- Useful when you do not have control over the thread creation process (i.e., when using a thread pool)
- Different from local variables
 - Local variables visible only during single function invocation
 - TLS visible across function invocations
- Similar to `static` data
 - TLS is unique to each thread

Threading Design Issue: Scheduler Activations

- Problem: Both M:M and Two-level models require communication to maintain the appropriate number of kernel threads allocated to the application
- Solution: use Light Weight Process (LWP), an intermediate data structure between user and kernel threads
 - Appears to be a virtual processor on which process can schedule user thread to run
 - Each LWP attached to kernel thread
 - Kernel raises upcalls to user-level threading library to give the information on kernel threads for the threading library to maintain the correct number kernel threads

