


ITP30002 Operating System

Process

This lecture note is taken from the instructor's resource of *Operating System Concept*, 10/e and then partly edited/revised by Shin Hong.

Need of Multiprogramming



- Multiprogramming: execute multiple programs on a single computer at the same time
 - Time-sharing: run multiple programs concurrently by dispatching each program execution to a CPU for a short time period
- Benefits
 - Improve throughput
 - Improve responsiveness (timesharing)

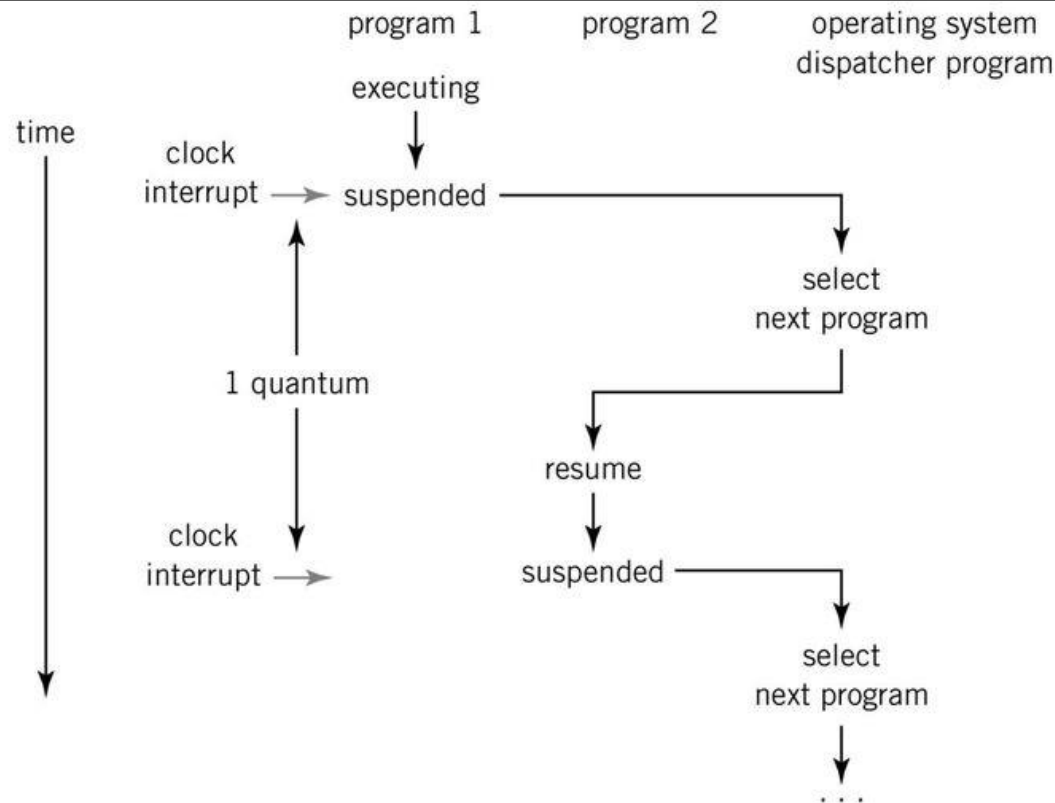
What is Process

- A **process** is a running instance of a program
 - a process is an object of OS to represent a status of a program execution
 - a process progresses in an instruction-by-instruction fashion
 - a program is a *passive* entity resided in a storage (e.g. an executable file) whereas a process is an *active* entity of the CPU and memory status
 - a process starts when executable file loaded into memory
 - multiple users may run multiple processes for a single program

Process Virtualizes Computer

- A **process** is an object to hold all entities to represent the running status of a program
 - CPU snapshot
 - Program counter
 - Registers
 - Memory snapshot
 - program code (also called text section)
 - Stack section: Function parameters, return addresses, local variables
 - Data section: memory for containing global variables
 - Heap section: memory for dynamically allocated during run time
 - State
- An OS can make multiple processes run at the same time by time-sharing a processor and space-sharing main memory

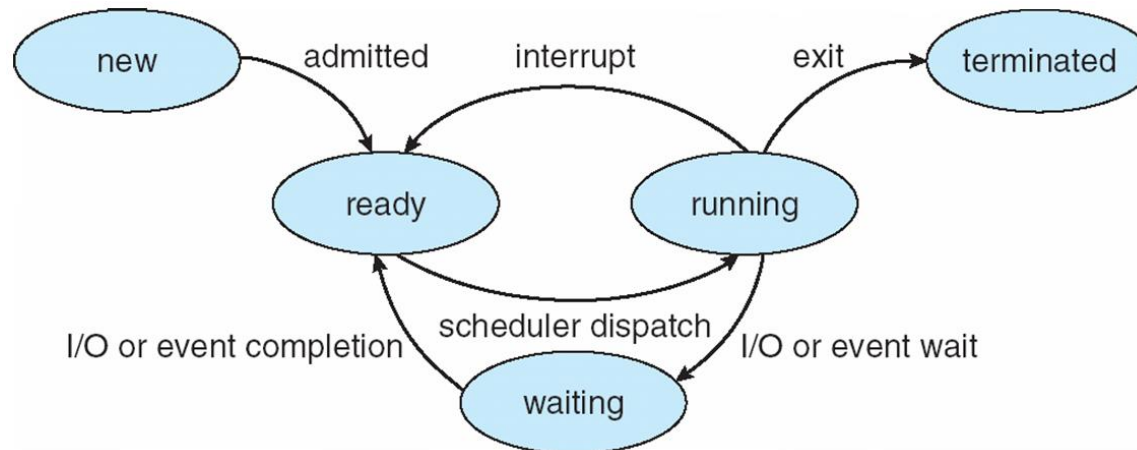
Time Sharing of Multiple Processes




Englander: The Architecture of Computer
Hardware and Systems Software, 2nd edition
Chapter 8, Figure 08-08

Process State

- As a process executes, the process state changes
 - **new**: The process is being created
 - **ready**: The process is waiting to be assigned to a processor
 - **running**: Instructions are being executed
 - **waiting**: The process is waiting for some event to occur
 - **terminated**: The process has finished execution





Memory Structure of Process



- Text: immutable data
 - instructions (code)
 - constants
- Data: always allocated variables
 - global variables
- Stack: dynamically allocated variables for function calls
 - local variables
- Heap: arbitrary dynamically allocated variables

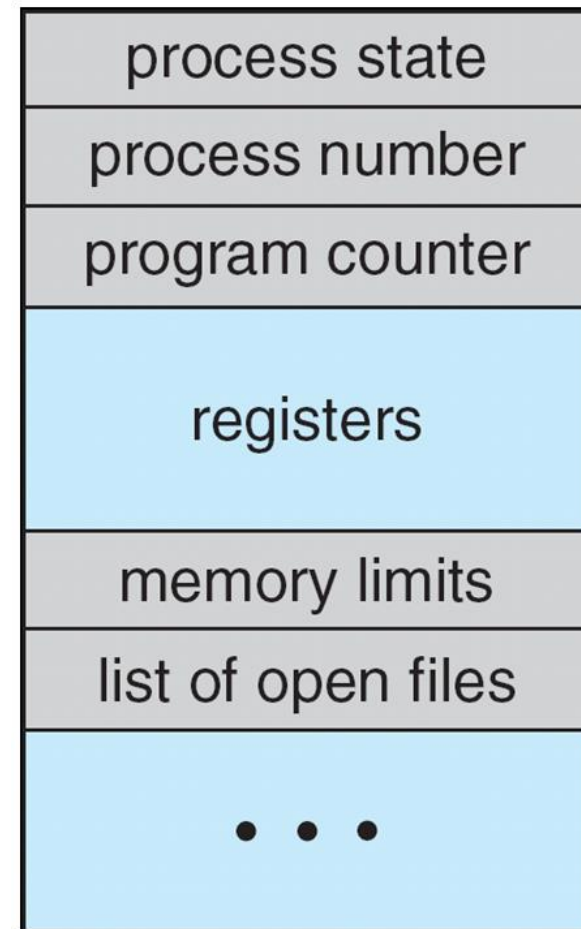
Process in Memory



Process Control Block (PCB)

Information associated with each process
(also called task control block)

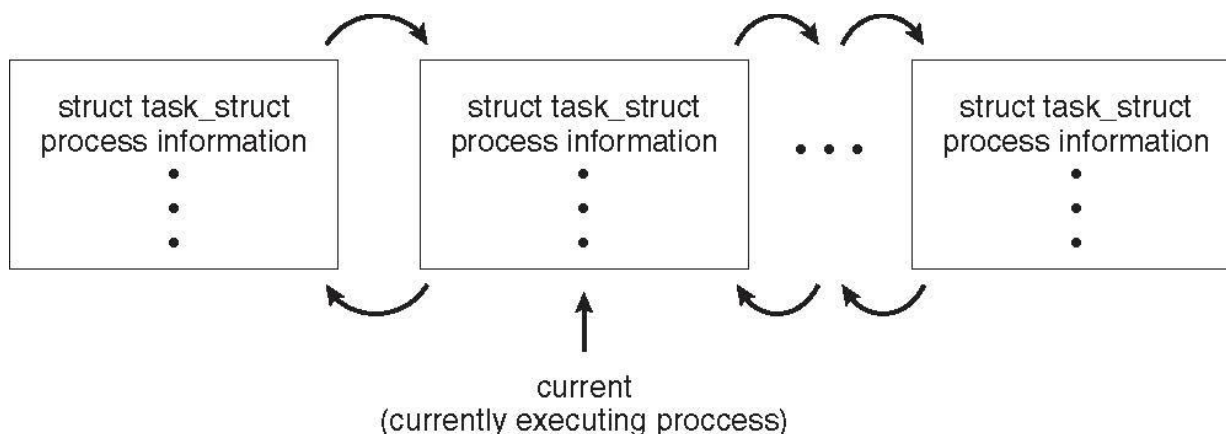
- Process state
- Program counter: a location of instruction to next execute
- CPU registers: contents of all process-centric registers
- CPU scheduling information: priorities, scheduling queue pointers
- Accounting information: CPU used, clock time elapsed since start, time limits
- Memory-management information: memory allocated to the process
- I/O status information: I/O devices allocated to process, list of open files



Process Representation in Linux

Represented by the C structure `task_struct`


| | | |
|-----------------------------------|-------------------------|------------------------------------------------|
| <code>pid</code> | <code>t_pid;</code> | <code>/* process identifier */</code> |
| <code>long</code> | <code>state;</code> | <code>/* state of the process */</code> |
| <code>unsigned int</code> | <code>time_slice</code> | <code>/* scheduling info. */</code> |
| <code>struct files_struct*</code> | <code>files;</code> | <code>/* list of open files */</code> |
| <code>struct mm_struct *</code> | <code>mm;</code> | <code>/* addr. space of this process */</code> |
| <code>struct task_struct *</code> | <code>parent;</code> | <code>/* this process's parent */</code> |
| <code>struct list_head</code> | <code>children;</code> | <code>/* this process's children */</code> |



Context Switch

- Context-switch is an action by OS that switches the running process to another one in the ready queue
 - save the state of the current process
 - select one of the ready processes
 - load the saved state for the selected process
- Context of a process represented in the PCB
- Context-switch time is overhead; the system does no useful work while switching
 - The more complex the OS and the PCB, the longer the context switch
 - Some hardware provides multiple sets of registers per CPU to load multiple contexts at once


Context Switch: Dispatching PCB to CPU



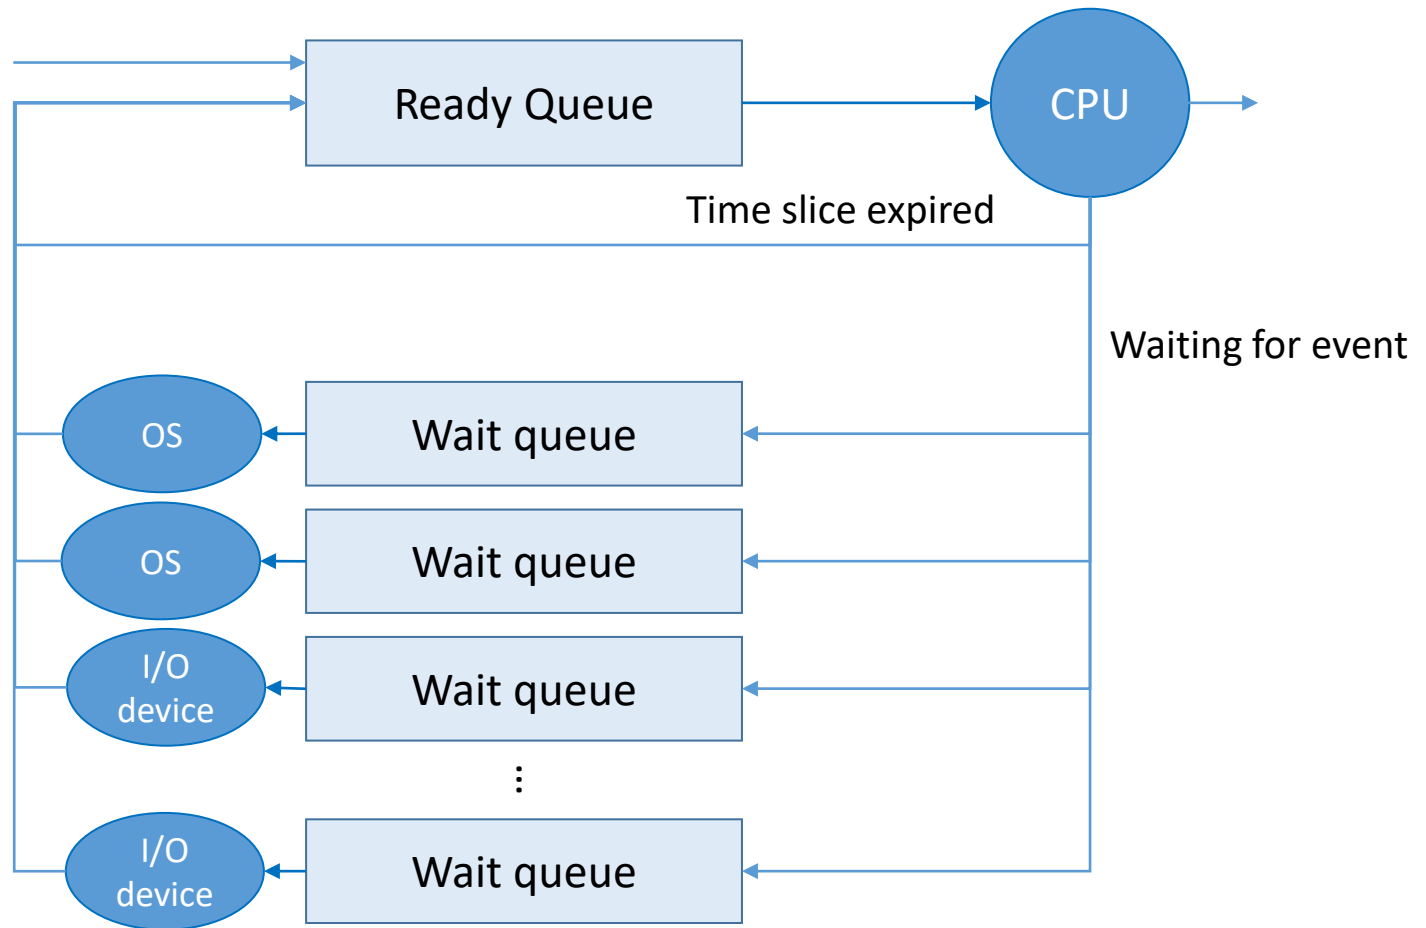
Process Scheduling

- Process scheduler selects one of available processes (ready state) for next execution on CPU
 - a scheduler is object to maximize CPU utilization and/or minimize response latency
- OS maintains scheduling queues of processes
 - Job queue: a set of all processes in the system
 - Ready queue: a set of all processes residing in main memory, ready and waiting to execute
 - Wait queues (device queues): a set of processes waiting for an I/O device (sleeping until an I/O completion)

Ready Queue And Various I/O Device Queues



Scheduling Lifecycle of Process




Schedulers

- **Short-term scheduler** (or CPU scheduler): selects which process should be executed next and allocates CPU
 - Sometimes the only scheduler in a system
 - Short-term scheduler is invoked frequently (milliseconds)
- **Long-term scheduler** (or job scheduler): selects which processes should be brought into the ready queue
 - Long-term scheduler is invoked infrequently (seconds, minutes)
 - The long-term scheduler controls the degree of multiprogramming

Addition of Medium Term Scheduling

- **Medium-term scheduler** can be added if a degree of multiple programming needs to decrease
 - **Swapping:** Remove a process from memory; store on disk; bring back in from disk to continue execution



Schedulers

- Processes can be described as either:
 - I/O-bound process: spends more time doing I/O than computations, many short CPU bursts
 - CPU-bound process: spends more time doing computations; few very long CPU bursts

Process vs. Threads

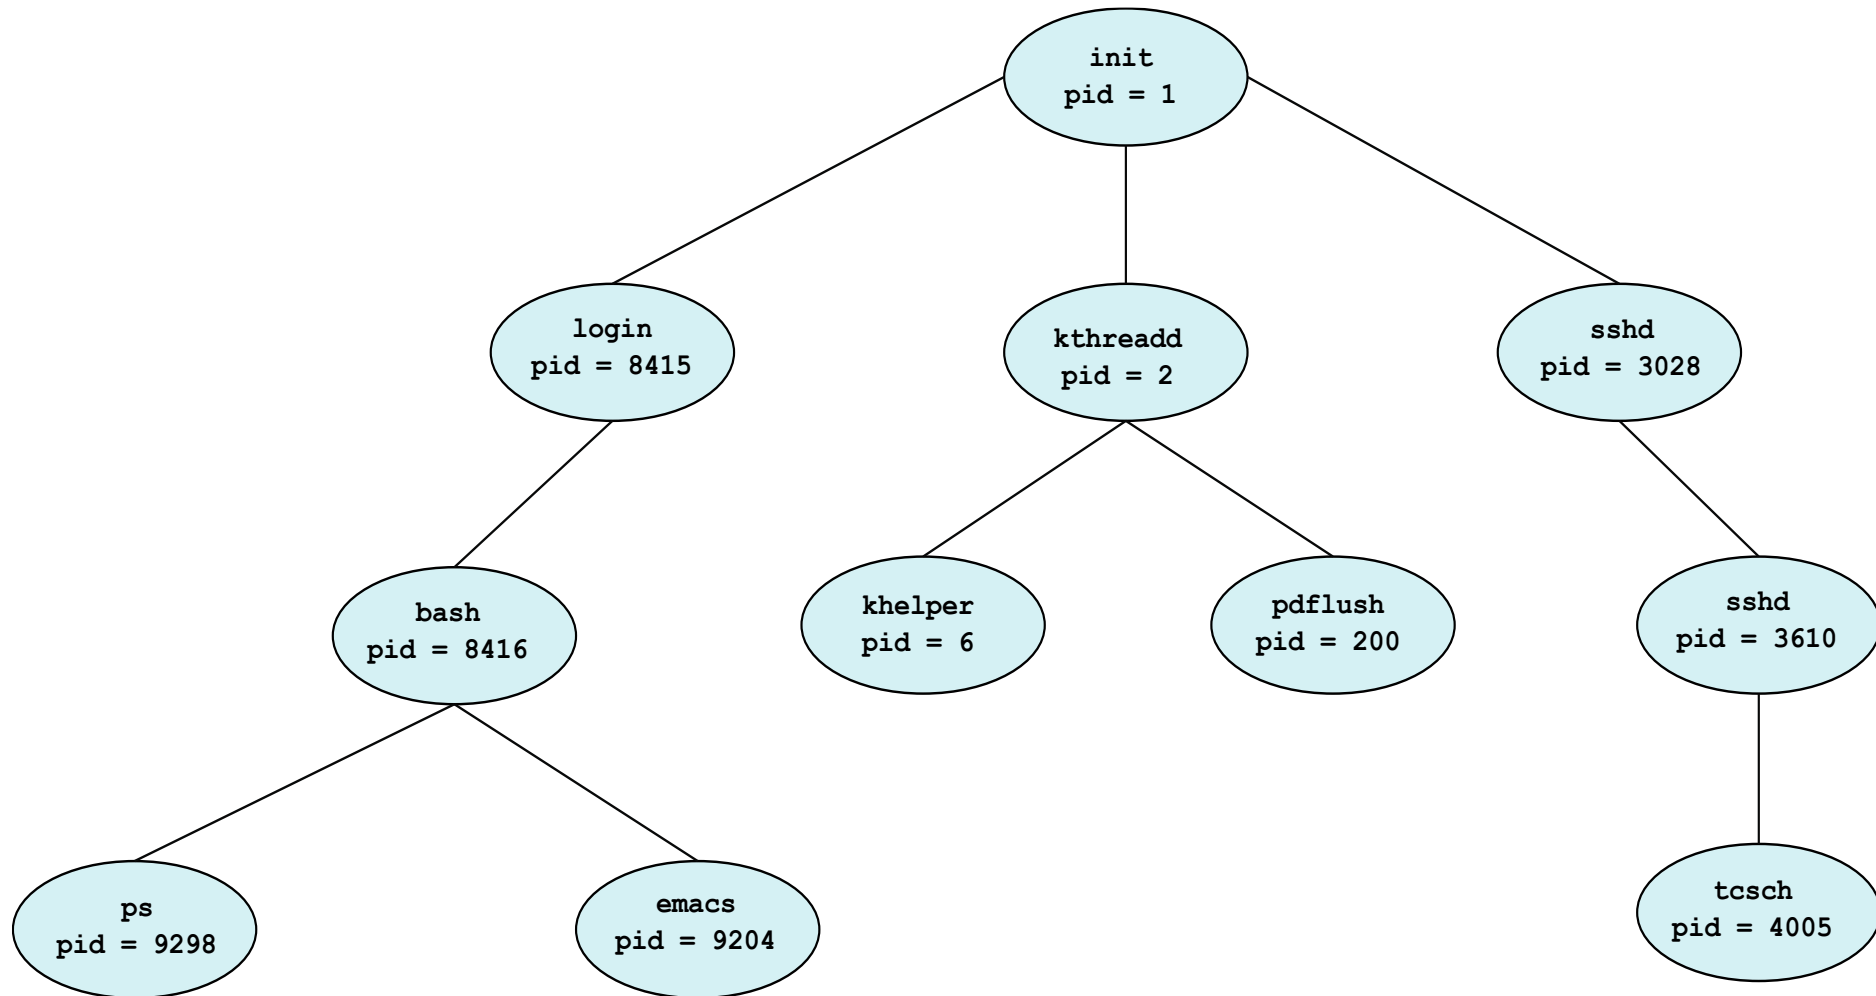
- A thread is an object to represent a single control flow
- A process with a single thread has a single flow of execution
- A process can have multiple threads
 - having multiple execution flows is called *multi-threading*
 - a PCB contains multiple entities for representing execution status (e.g., PC), meanwhile a PCB contains one memory status



Process Creation

- A process is identified and managed via a process identifier (pid)
- A parent process can spawn a child process to delegate a subtask
 - A process can spawn multiple children processes
 - A parent process can run concurrently with its children processes
 - A child process, in turn create other processes, forming a tree of processes
 - A parent can wait until a child (or children) terminates
- A parent and its children can share resources
 - Children may share a subset of parent's resources
- Process in UNIX
 - a system call **fork()** system call creates a new process
 - a child process duplicates the memory of its parent

A Tree of Processes in Linux



Process Termination

- Process executes last statement and then asks the operating system to delete it using the **exit()** system call
 - Returns status data from child to parent (via **wait()**)
 - Process' resources are deallocated by operating system
- Parent may terminate the execution of children processes using the **abort()** system call. Some reasons for doing so:
 - Child has exceeded allocated resources
 - Task assigned to child is no longer required
 - The parent is exiting and the operating systems does not allow a child to continue if its parent terminates

Process Termination

- Some OSes do not allow a child to exist if its parent has terminated. If a process terminates, then all its children must also be terminated.
 - Cascading termination; All children including grandchildren are terminated.
 - The termination is initiated by the operating system.
- The parent process may wait for termination of a child process by using the `wait()` system call. The call returns status information and the pid of the terminated process
 - A process is a zombie if no parent is waiting
 - A process is an orphan if its parent terminated (without invoking `wait()`)

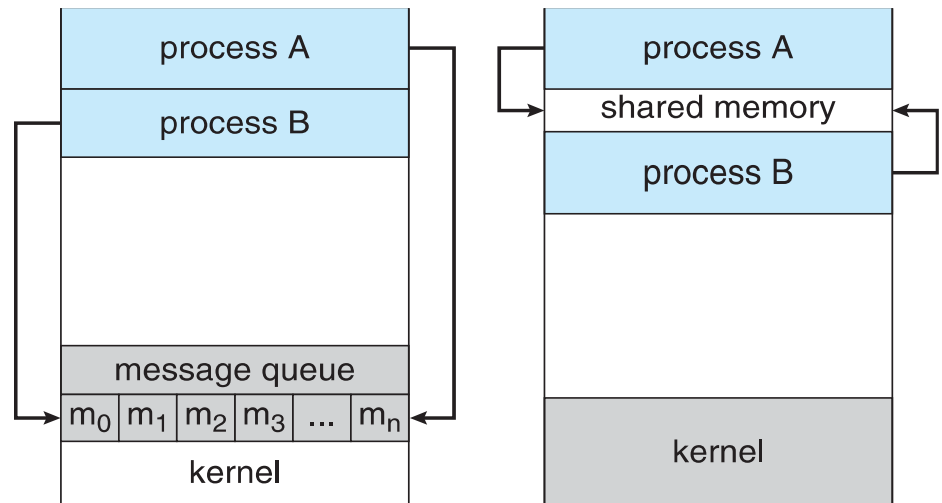
Multiprocess Architecture – Chrome Browser

- Many web browsers ran as a single process (some still do)
 - If one web site causes trouble, entire browser can hang or crash
- Google Chrome is multiprocess with 3 different processes:
 - Browser process manages user interface, disk and network I/O
 - Renderer process renders web pages, deals with HTML and Javascript.
A new renderer created for each website opened
 - Runs in sandbox restricting disk and network I/O, minimizing effect of security exploits
 - Plug-in process for each type of plug-in



Interprocess Communication

- Processes within a system may be **independent** or **cooperating**
 - Cooperating process can affect or be affected by other processes, including sharing data
- Cooperating processes need **inter-process communication (IPC)**
 - Shared memory**
 - Message passing**



(a) Message passing. (b) shared memory.

Message Passing

- Mechanism for processes to communicate with each other without resorting to shared variables
 - If processes P and Q wish to communicate, they need to:
 - Establish a **communication link** between them
 - Exchange messages via send/receive
- IPC facility provides two operations:
 - **send**(message)
 - **receive**(message)
- Implementation issue
 - Physical media: shared memory, Hardware bus, Network
 - Direct or indirect
 - Synchronous or asynchronous
 - Automatic or explicit buffering

Direct Communication

- Processes must name each other explicitly:
 - **send** ($P, message$) – send a message to process P
 - **receive**($Q, message$) – receive a message from process Q
- Properties of communication link
 - Links are established automatically
 - A link is associated with exactly one pair of communicating processes
 - Between each pair there exists exactly one link
 - The link may be unidirectional, but is usually bi-directional

Indirect Communication

- Messages are directed and received from mailboxes (also referred to as ports)
 - Each mailbox has a unique ID
 - Processes can communicate only if they share a mailbox
 - Primitives are defined as:
 - `send(A, message)` – send a message to mailbox A
 - `receive(A, message)` – receive a message from mailbox A
- Properties of communication link
 - Link established only if processes share a common mailbox
 - A link may be associated with many processes
 - Each pair of processes may share several communication links
 - Link may be unidirectional or bi-directional

Synchronization

- Message passing may be either blocking or non-blocking
- Blocking is considered synchronous
 - the sender is blocked until the message is received
 - the receiver is blocked until a message is available
- Non-blocking is considered asynchronous
 - the sender sends the message and continue
 - the receiver receives a valid message, or null message
- Different combinations possible
 - If both send and receive are blocking, we have a rendezvous

Buffering

- Queue of messages attached to the link.
- implemented in one of three ways
 1. Zero capacity – no messages are queued on a link.
Sender must wait for receiver (rendezvous)
 2. Bounded capacity – finite length of n messages
Sender must wait if link full
 3. Unbounded capacity – infinite length
Sender never waits

Examples of IPC Systems - POSIX

- How to use POSIX Shared Memory

- Process first creates/opens shared memory segment as a file

```
shm_fd = shm_open(name, O_CREAT | O_RDWR, 0666);
```

- Set the size of the object

```
ftruncate(shm_fd, 4096);
```

- Map the shared memory file to a memory

```
ptr = mmap(0, 4096, PROT_WRITE, M_SHARED, shm_fd, 0);
```

- Now the process could write to the shared memory

```
sprintf(ptr, "Writing to shared memory");
```

Producer-Consumer Example

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <fcntl.h>
#include <sys/shm.h>
#include <sys/stat.h>

int main()
{
    /* the size (in bytes) of shared memory object */
    const int SIZE = 4096;
    /* name of the shared memory object */
    const char *name = "OS";
    /* strings written to shared memory */
    const char *message_0 = "Hello";
    const char *message_1 = "World!";

    /* shared memory file descriptor */
    int shm_fd;
    /* pointer to shared memory object */
    void *ptr;

    /* create the shared memory object */
    shm_fd = shm_open(name, O_CREAT | O_RDWR, 0666);

    /* configure the size of the shared memory object */
    ftruncate(shm_fd, SIZE);

    /* memory map the shared memory object */
    ptr = mmap(0, SIZE, PROT_WRITE, MAP_SHARED, shm_fd, 0);

    /* write to the shared memory object */
    sprintf(ptr, "%s", message_0);
    ptr += strlen(message_0);
    sprintf(ptr, "%s", message_1);
    ptr += strlen(message_1);

    return 0;
}
```

```
#include <stdio.h>
#include <stdlib.h>
#include <fcntl.h>
#include <sys/shm.h>
#include <sys/stat.h>

int main()
{
    /* the size (in bytes) of shared memory object */
    const int SIZE = 4096;
    /* name of the shared memory object */
    const char *name = "OS";
    /* shared memory file descriptor */
    int shm_fd;
    /* pointer to shared memory object */
    void *ptr;

    /* open the shared memory object */
    shm_fd = shm_open(name, O_RDONLY, 0666);

    /* memory map the shared memory object */
    ptr = mmap(0, SIZE, PROT_READ, MAP_SHARED, shm_fd, 0);

    /* read from the shared memory object */
    printf("%s", (char *)ptr);

    /* remove the shared memory object */
    shm_unlink(name);

    return 0;
}
```

Pipes

- Acts as a channel allowing two processes to communicate
 - **Ordinary pipes** are typically created by a parent process to communicate with a child process; they cannot be accessed from outside the process that created it.
 - **Named pipes** can be accessed without a parent-child relationship
- Ex. Ordinary pipe

```
#include <sys/types.h>
#include <stdio.h>
#include <string.h>
#include <unistd.h>

#define BUFFER_SIZE 25
#define READ_END 0
#define WRITE_END 1

int main(void)
{
    char write_msg[BUFFER_SIZE] = "Greetings";
    char read_msg[BUFFER_SIZE];
    int fd[2];
    pid_t pid;

    /* create the pipe */
    if (pipe(fd) == -1) {
        fprintf(stderr, "Pipe failed");
        return 1;
    }
}
```

```
/* fork a child process */
pid = fork();

if (pid > 0) { /* parent process */
    /* close the unused end of the pipe */
    close(fd[READ_END]);

    /* write to the pipe */
    write(fd[WRITE_END], write_msg, strlen(write_msg)+1);

    /* close the write end of the pipe */
    close(fd[WRITE_END]);
}
else { /* child process */
    /* close the unused end of the pipe */
    close(fd[WRITE_END]);

    /* read from the pipe */
    read(fd[READ_END], read_msg, BUFFER_SIZE);
    printf("read %s", read_msg);

    /* close the write end of the pipe */
    close(fd[READ_END]);
}
```

Example

```
#include <sys/types.h>
#include <stdio.h>
#include <string.h>
#include <unistd.h>

#define BUFFER_SIZE 25
#define READ_END 0
#define WRITE_END 1

int main(void)
{
    char write_msg[BUFFER_SIZE] = "Greetings";
    char read_msg[BUFFER_SIZE];
    int fd[2];
    pid_t pid;

    /* create the pipe */
    if (pipe(fd) == -1) {
        fprintf(stderr, "Pipe failed");
        return 1;
    }

    /* create the pipe */
    if (pipe(fd) == -1) {
        fprintf(stderr, "Pipe failed");
        return 1;
    }

    /* fork a child process */
    pid = fork();

    if (pid > 0) { /* parent process */
        /* close the unused end of the pipe */
        close(fd[READ_END]);

        /* write to the pipe */
        write(fd[WRITE_END], write_msg, strlen(write_msg)+1);


        /* close the write end of the pipe */
        close(fd[WRITE_END]);
    }
    else { /* child process */
        /* close the unused end of the pipe */
        close(fd[WRITE_END]);

        /* read from the pipe */
        read(fd[READ_END], read_msg, BUFFER_SIZE);
        printf("read %s", read_msg);

        /* close the write end of the pipe */
        close(fd[READ_END]);
    }
}
```

Sockets

- A **socket** is defined as an endpoint for communication
- Concatenation of IP address and **port** – a number included at start of message packet to differentiate network services on a host
- The socket **161.25.19.8:1625** refers to port **1625** on host **161.25.19.8**
- Communication consists between a pair of sockets
- All ports below 1024 are **well known**, used for standard services



Remote Procedure Calls (RPC)

- RPC abstracts procedure calls between processes on networked systems
 - Again uses ports for service differentiation
- Stubs – client-side proxy for the actual procedure on the server
- The client-side stub locates the server and marshalls the parameters
 - Data representation handled via External Data Representation (XDL) format to account for different architectures
 - e.g., Big-endian and little-endian
- The server-side stub receives this message, unpacks the marshalled parameters, and performs the procedure on the server, and then send back the output to the client
 - Remote communication has more failure scenarios than local
 - OS typically provides a rendezvous (or matchmaker) service to connect client and server

Execution of RPC

